


CSCI 544: Applied Natural Language Processing

# **Seq2seq Generation & Neural Machine Translation**

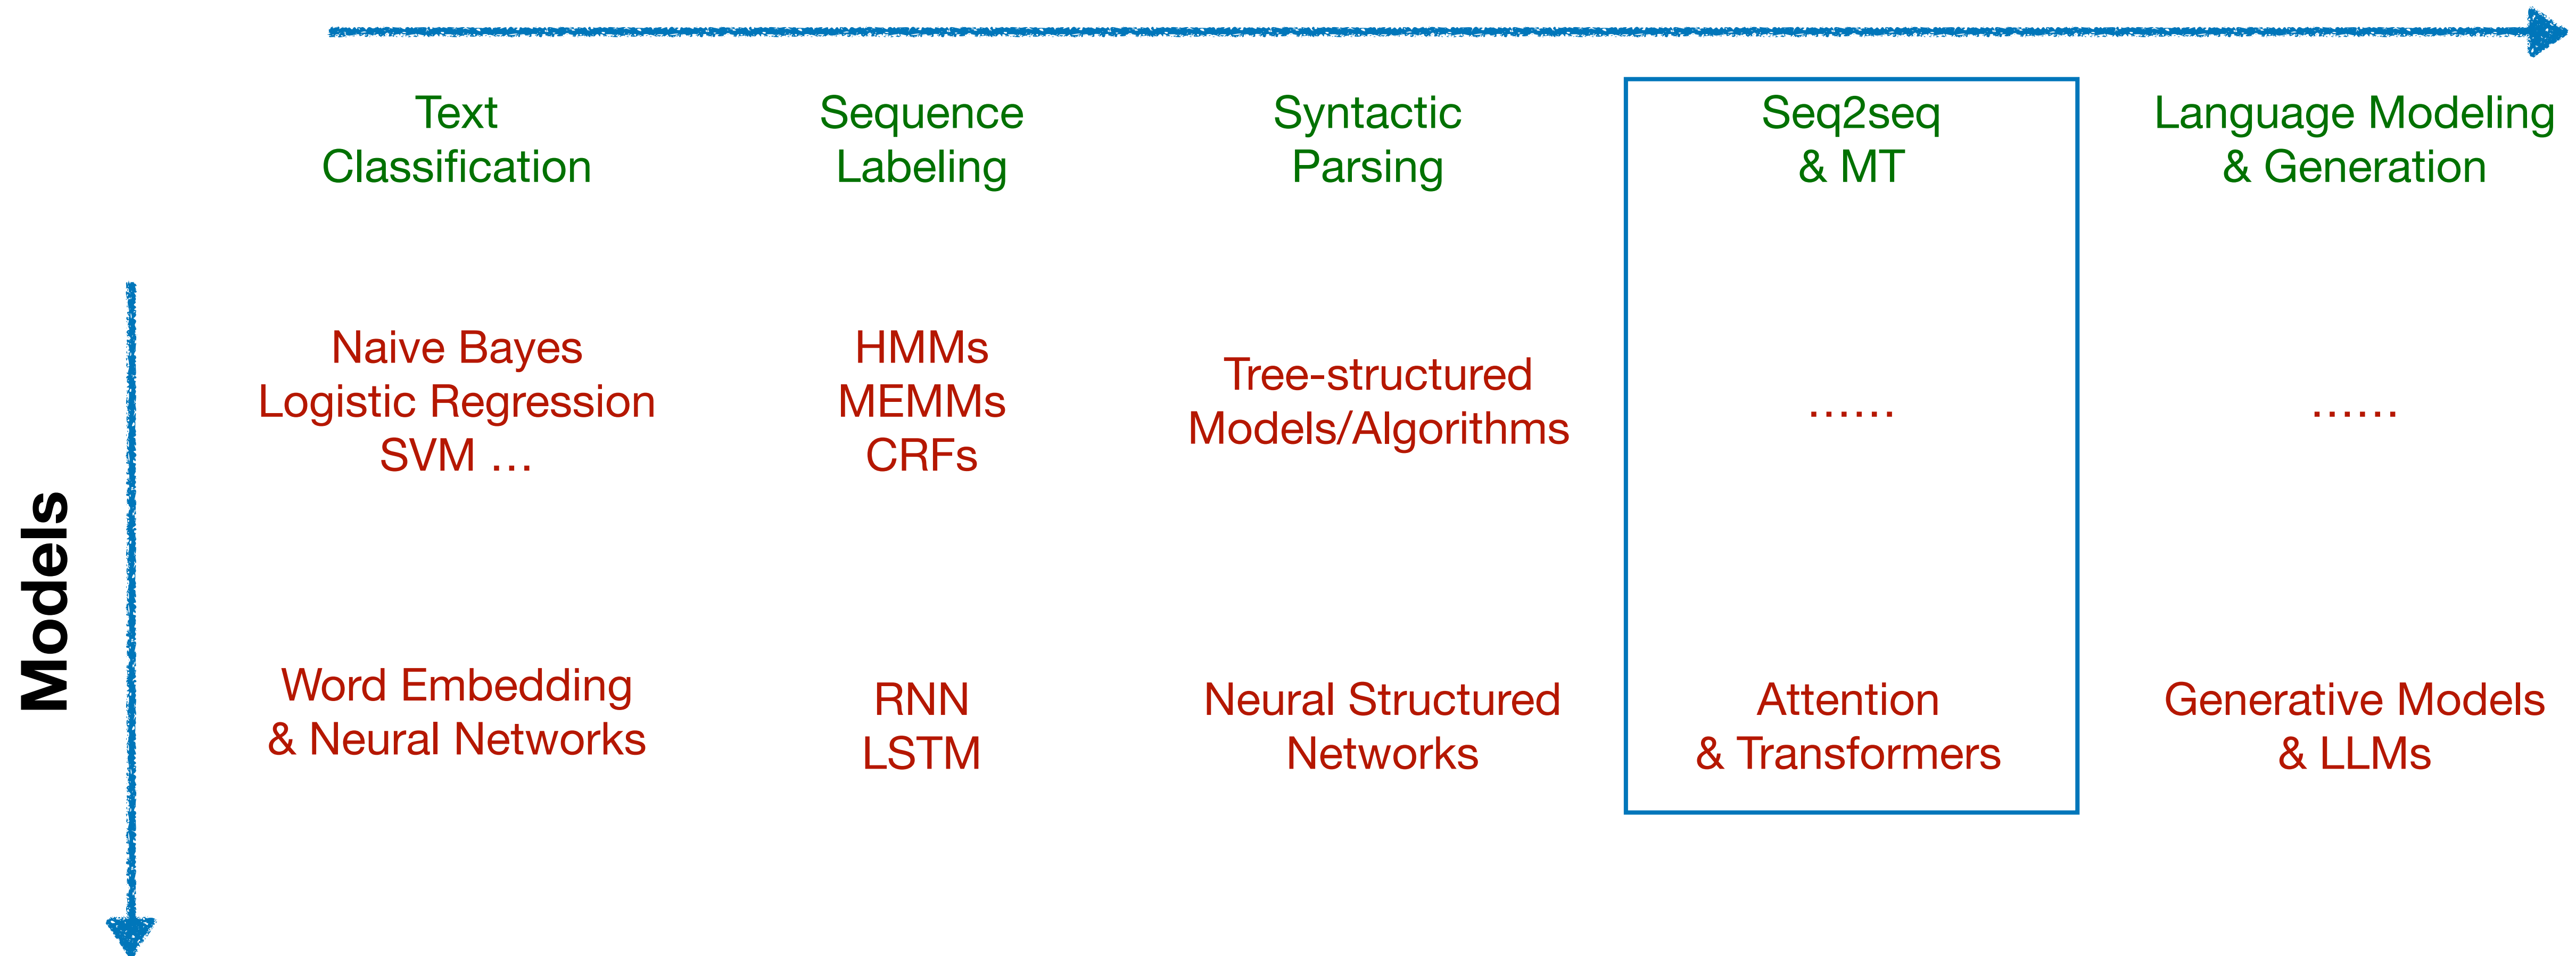
Xuezhe Ma (Max)



**USC** University of  
Southern California

# Course Organization


## NLP Tasks



# Seq2seq Generation

- Sequence-to-Sequence (Seq2seq) Generation

- Input:  $X = \{x_1, x_2, \dots, x_L\}, x_i \in \mathcal{X}$
- Output:  $Y = \{y_1, y_2, \dots, y_T\}, y_i \in \mathcal{Y}$
- Model:  $p_\theta(Y|X)$



# Seq2seq Generation


- **Sequence-to-Sequence (Seq2seq) Generation**

- Input:  $X = \{x_1, x_2, \dots, x_L\}, x_i \in \mathcal{X}$
- Output:  $Y = \{y_1, y_2, \dots, y_T\}, y_i \in \mathcal{Y}$
- Model:  $p_{\theta}(Y|X)$

<u>Input <math>X</math></u>	<u>Output <math>Y</math> (<b>Text</b>)</u>	<u>Task</u>
English	Japanese	Translation
Document	Short Description	Summarization
Utterance	Response	Response Generation
Image	Text	Image Captioning
Speech	Transcript	Speech Recognition

# Seq2seq Generation

- **Sequence-to-Sequence (Seq2seq) Generation**
  - Input:  $X = \{x_1, x_2, \dots, x_L\}, x_i \in \mathcal{X}$
  - Output:  $Y = \{y_1, y_2, \dots, y_T\}, y_i \in \mathcal{Y}$
  - Model:  $p_{\theta}(Y|X)$  **How?**
- **Difference from Sequence Labeling**
  - The length of  $Y$  can be different from the length of  $X$
  - The space of  $\mathcal{Y}$  is often much larger



# Statistic Machine Translation

# Statistical Machine Translation

- **IBM Translation Models**
  - Word-level alignment model
  - EM algorithm
- **Phrase-based Translation Models**
  - Phrase-based alignment model
- **Heavy Engineering**
  - Moses system
  - 360 pages manual

# Statistical Machine Translation

- **IBM Translation Models**
  - Word-level alignment model
  - EM algorithm
- **Phrase-based Translation Models**
  - Phrase-based alignment model
- **Heavy Engineering**
  - Moses system
  - 360 pages manual



# Word-Alignment Model in SMT

- **Key Idea:** two words are more likely to be aligned when they occur more frequently in translation pairs

我 不 知道

I don't know

我 是 学生


I am a student

我 爱 喝 茶

I love drinking tea

# Word-Alignment Model in SMT

- $e$  is an English sentence with  $l$  words
- $f$  is a foreign sentence with  $m$  words
- An alignment  $a = \{a_1, a_2, \dots, a_m\}$ ,  $a_j \in \{0, \dots, l\}$
- Hence there are  $(l + 1)^m$  possible alignments



# Word-Alignment Model in SMT

- IBM Model 1:

$$p(a | e, m) = \frac{1}{(l + 1)^m}$$

- IBM Model 2:

$$p(a | e, m) = q(a_j | j, l, m)$$

- IBM Model 3, 4, 5, 6...

Translation Model

$$p(f | e) = \sum_{a \in \mathcal{A}} p(a | e, m) p(f | a, e, m)$$


# Statistical Machine Translation

- **IBM Translation Models**
  - Word-level alignment model
  - EM algorithm
- **Phrase-based Translation Models**
  - Phrase-based alignment model
- **Heavy Engineering**
  - Moses system
  - 360 pages manual

# Neural Machine Translation

# Seq2seq Generation

- **Sequence-to-Sequence (Seq2seq) Generation**
  - Input:  $X = \{x_1, x_2, \dots, x_L\}, x_i \in \mathcal{X}$
  - Output:  $Y = \{y_1, y_2, \dots, y_T\}, y_i \in \mathcal{Y}$
  - Model:  $p_\theta(Y|X)$  **How?**
- **Difference from Sequence Labeling**
  - The length of  $Y$  can be different from the length of  $X$
  - The size of  $\mathcal{Y}$  is often much larger

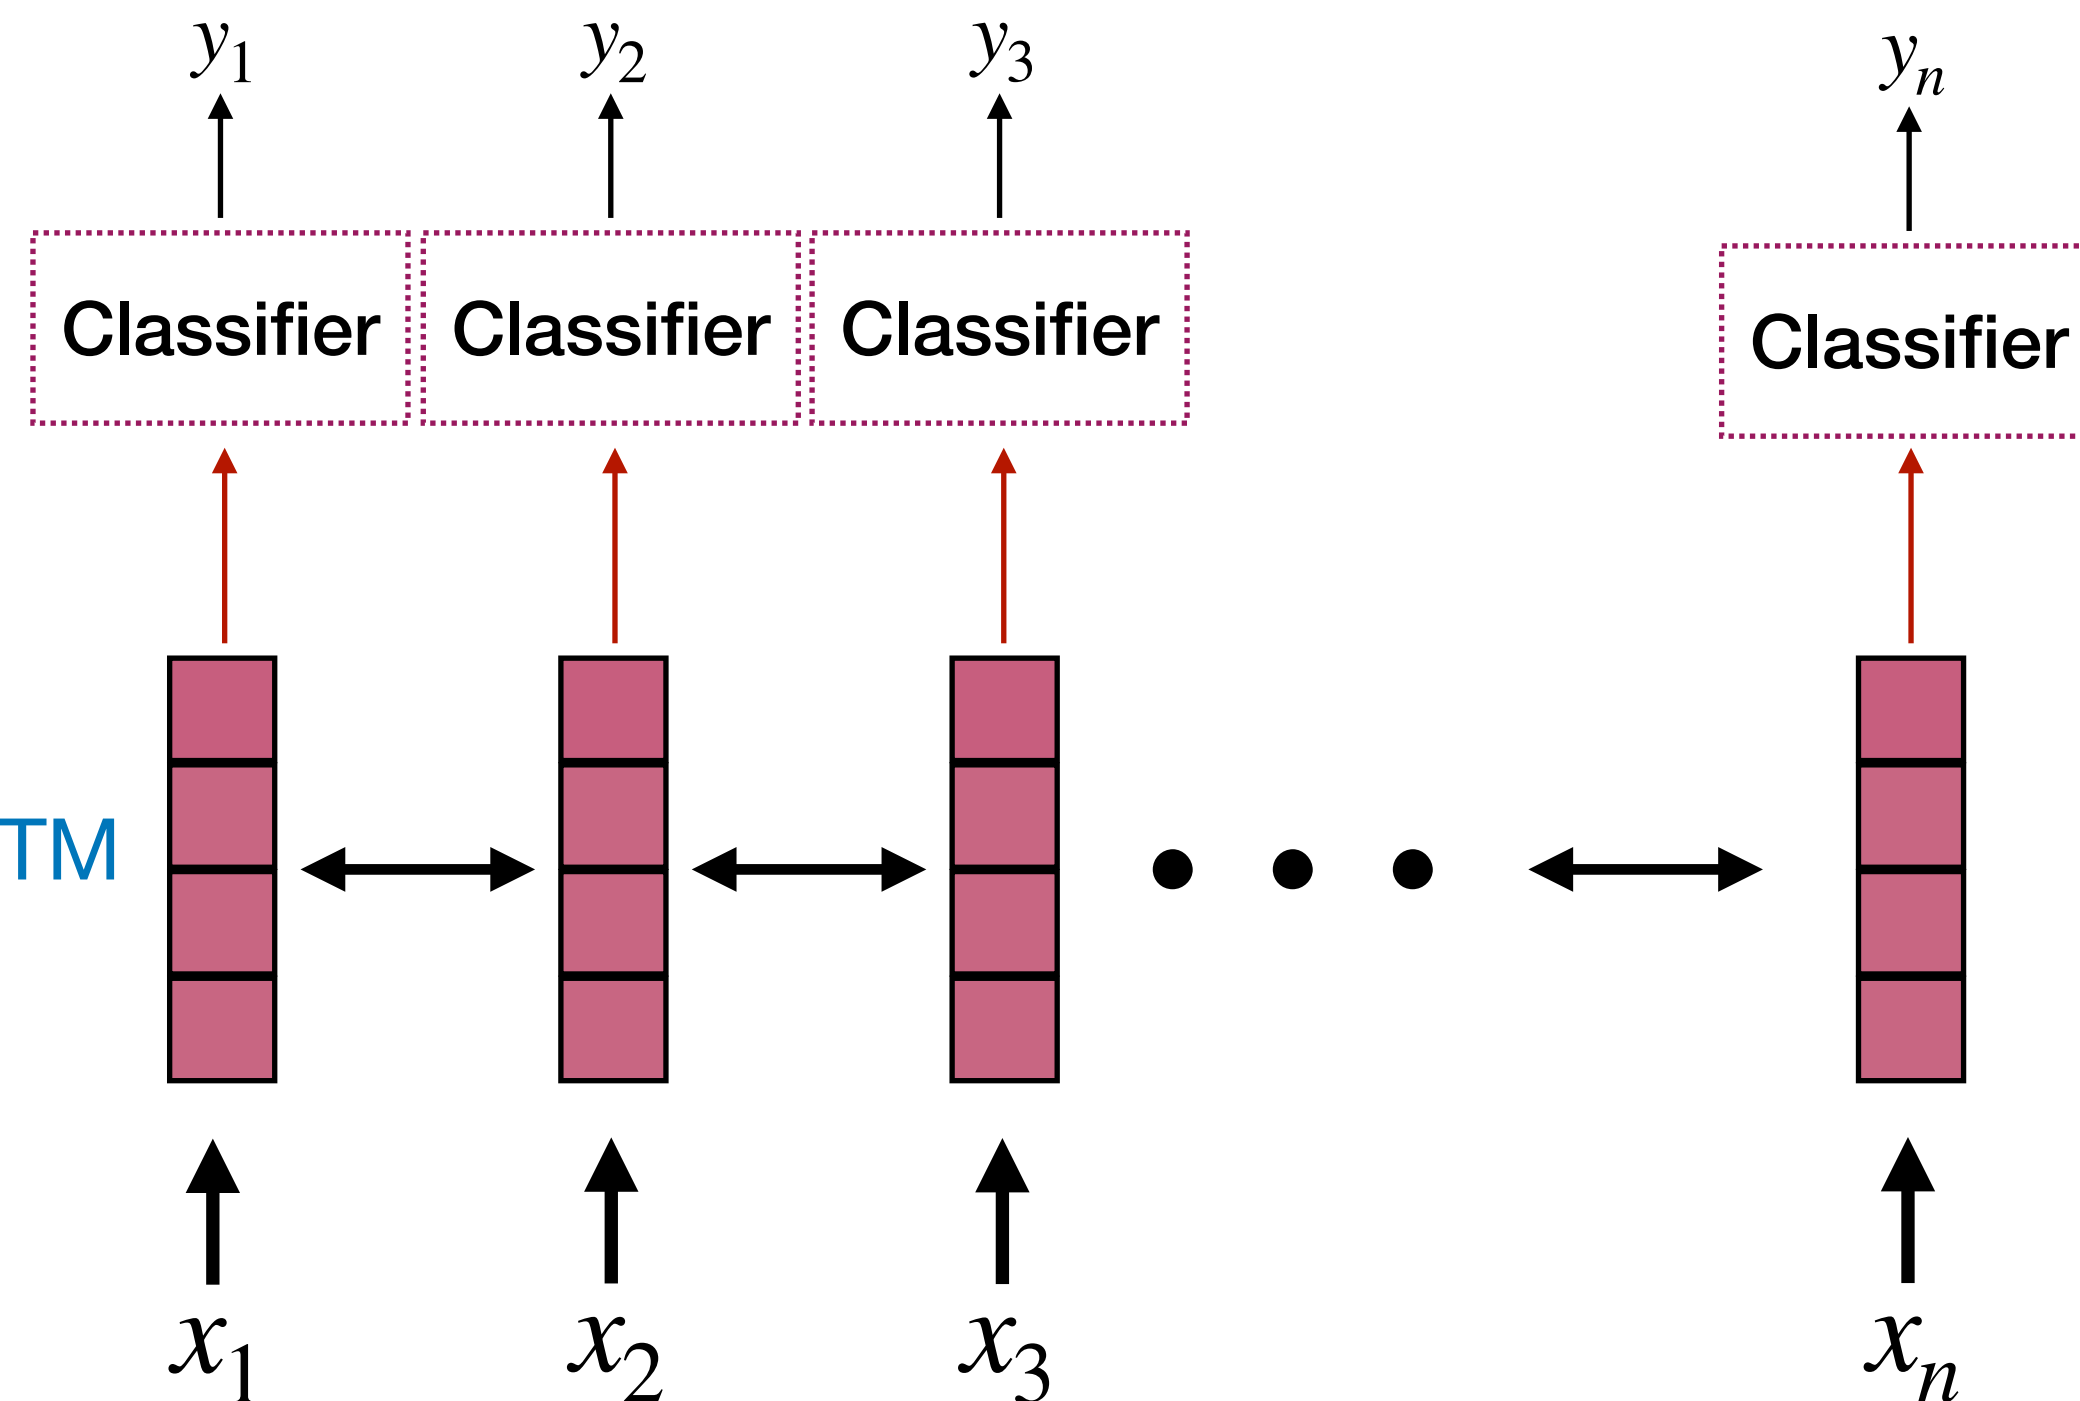


# Autoregressive Seq2seq Generation

- Sequence labeling vs. Seq2seq Generation

## Sequence labeling

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t|X)$$



Why not for seq2seq generation?

# Autoregressive Seq2seq Generation

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t|X)$$

Not a good choice!

我 不 知道

**I don't know**


**I do not know**

**I have no idea**



# Autoregressive Seq2seq Generation

- Autoregressive Factorization:

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$


- Autoregressive factorization is just chain-rule (HMMs, MEMMs)
- Autoregressive factorization does **NOT** assume any independence
- With autoregressive factorization, we need to model each  $p_{\theta}(y_t | y_{<t}, X)$




# Encoder-Decoder Architecture

- Sequence labeling vs. Seq2seq Generation

## Sequence labeling

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | \boxed{X})$$



## Seq2seq Generation


$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(\boxed{y_t} | y_{<t}, \boxed{X})$$

**Encoder:** encode a sentence into a sequence of vectors

**Decoder:** use another LSTM?


# Encoder-Decoder Architecture

- **Two Components:**
  - **Encoder:** Convert input sequence into a sequence of vectors
  - **Decoder:** Convert encoding into a sequence in the output space



# Special Tokens in Seq2seq

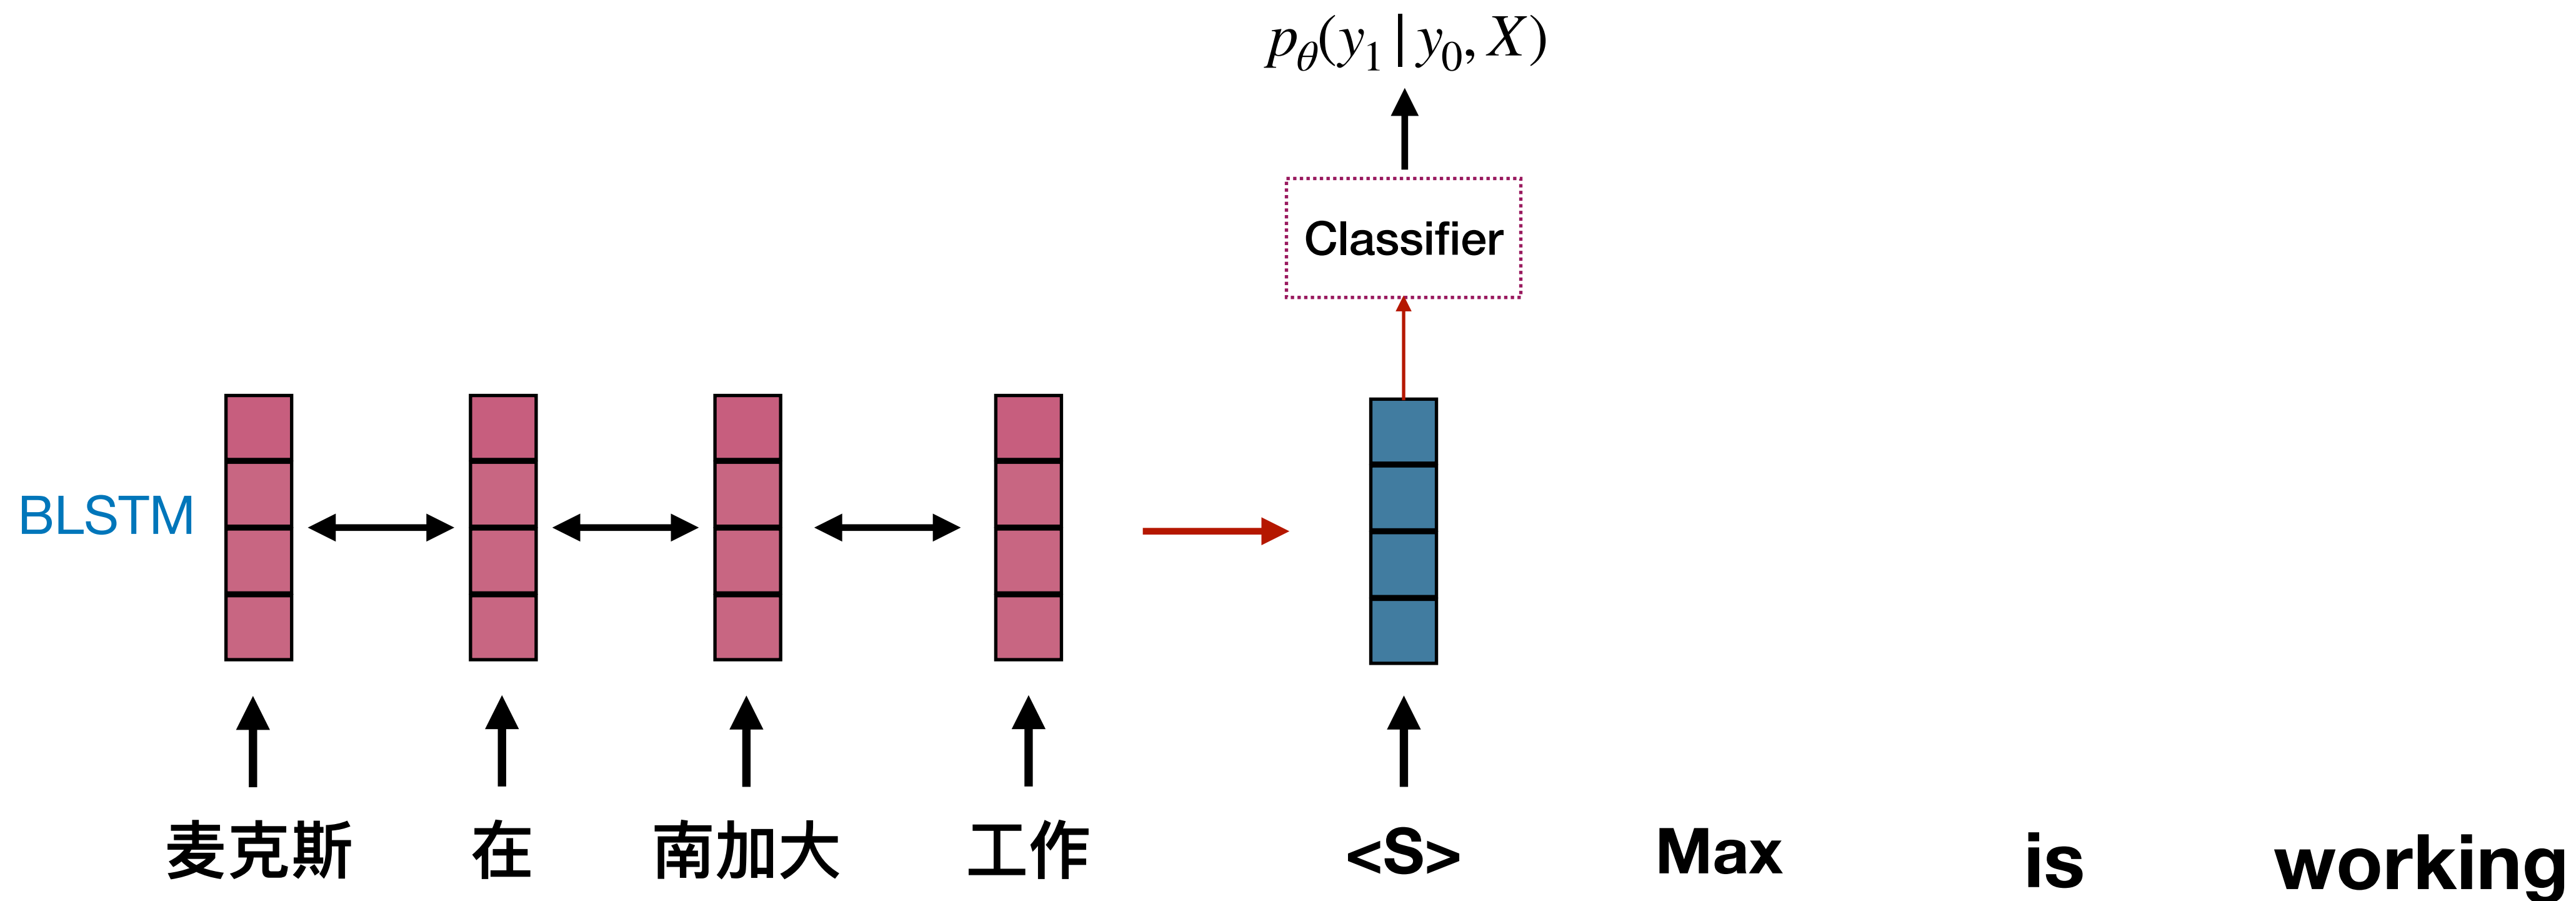
- **<BOS>**: start of the target sentence
- **<EOS>**: end of the target sentence



# Seq2seq Training

- Model Training:

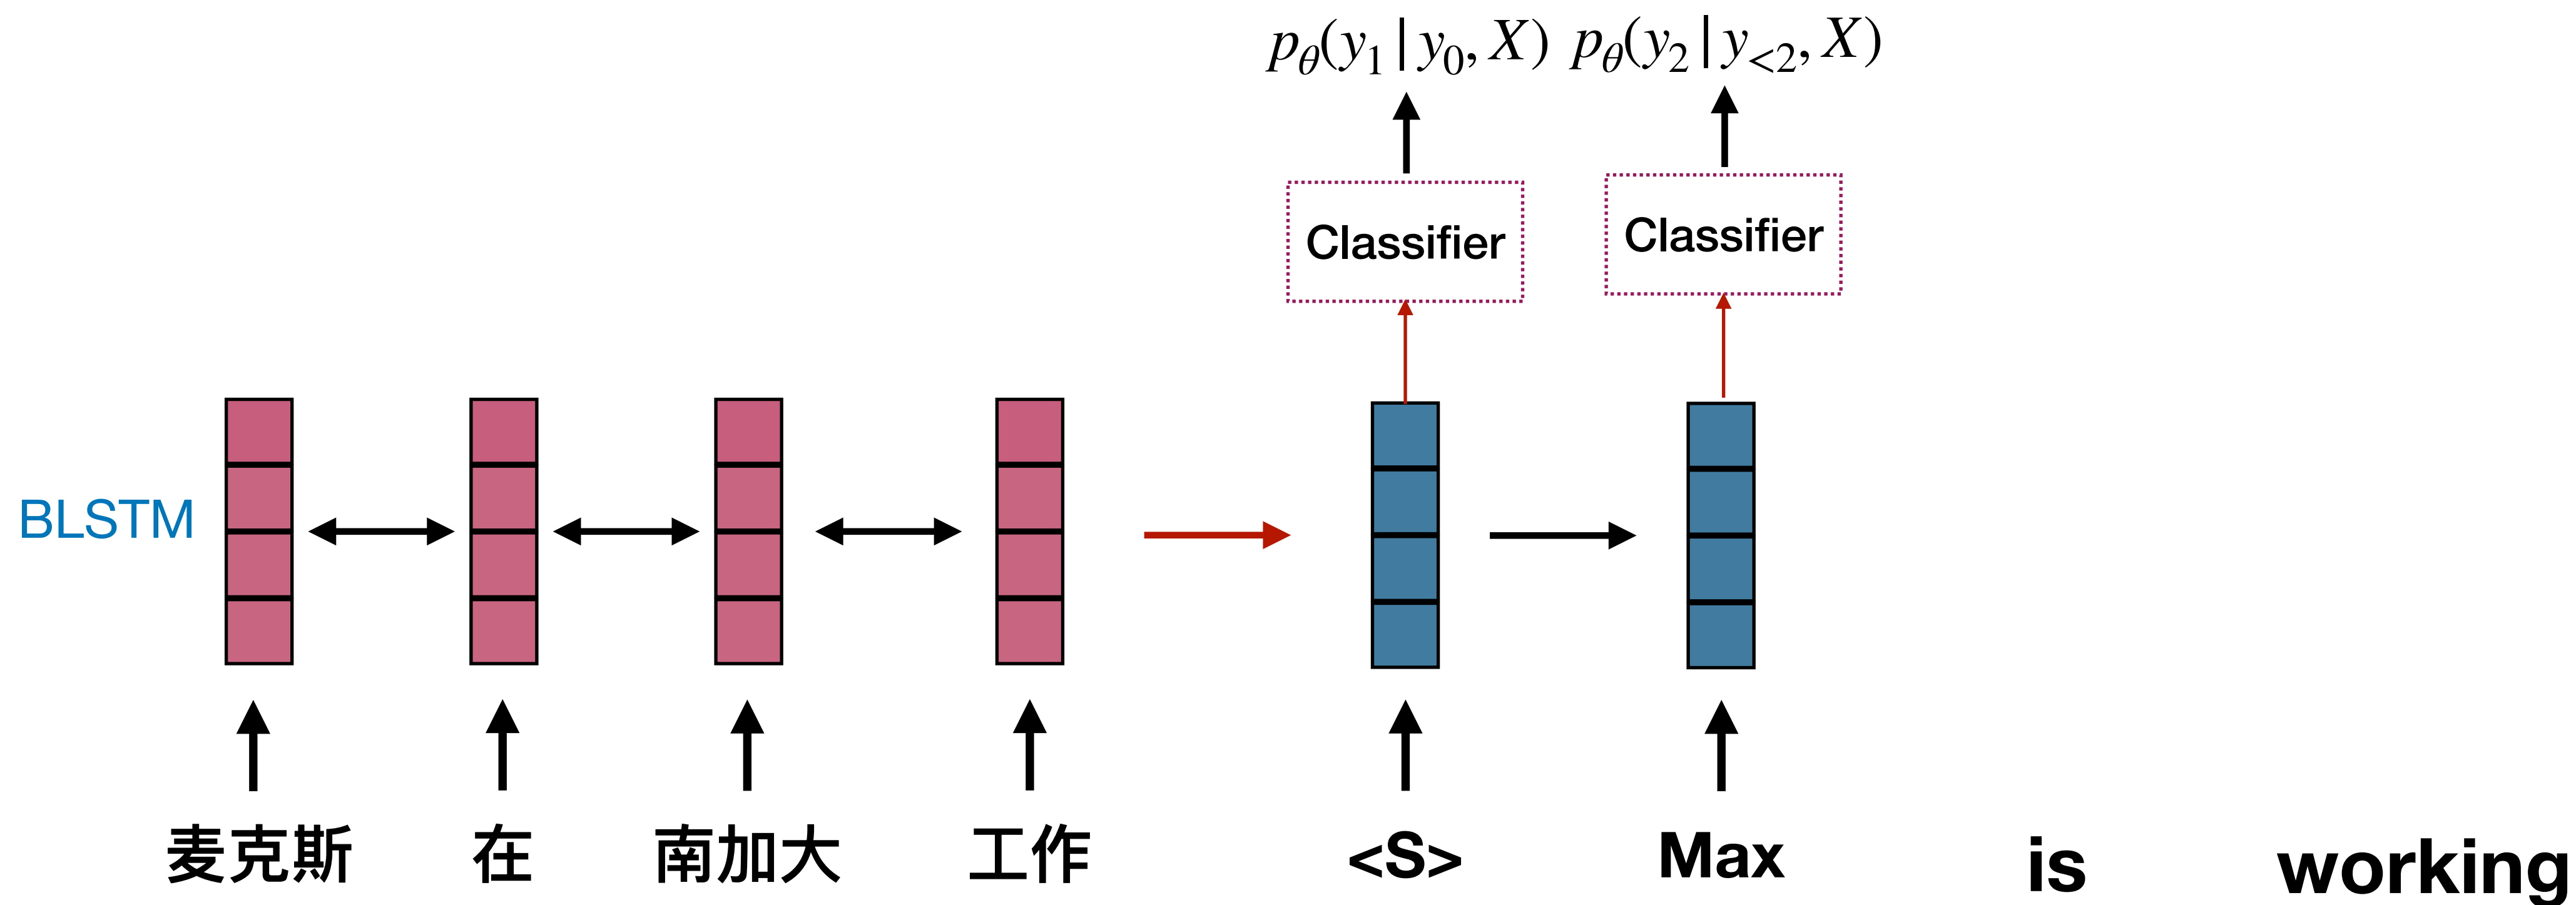
$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X) \quad t = 1$$



# Seq2seq Training

- Model Training:

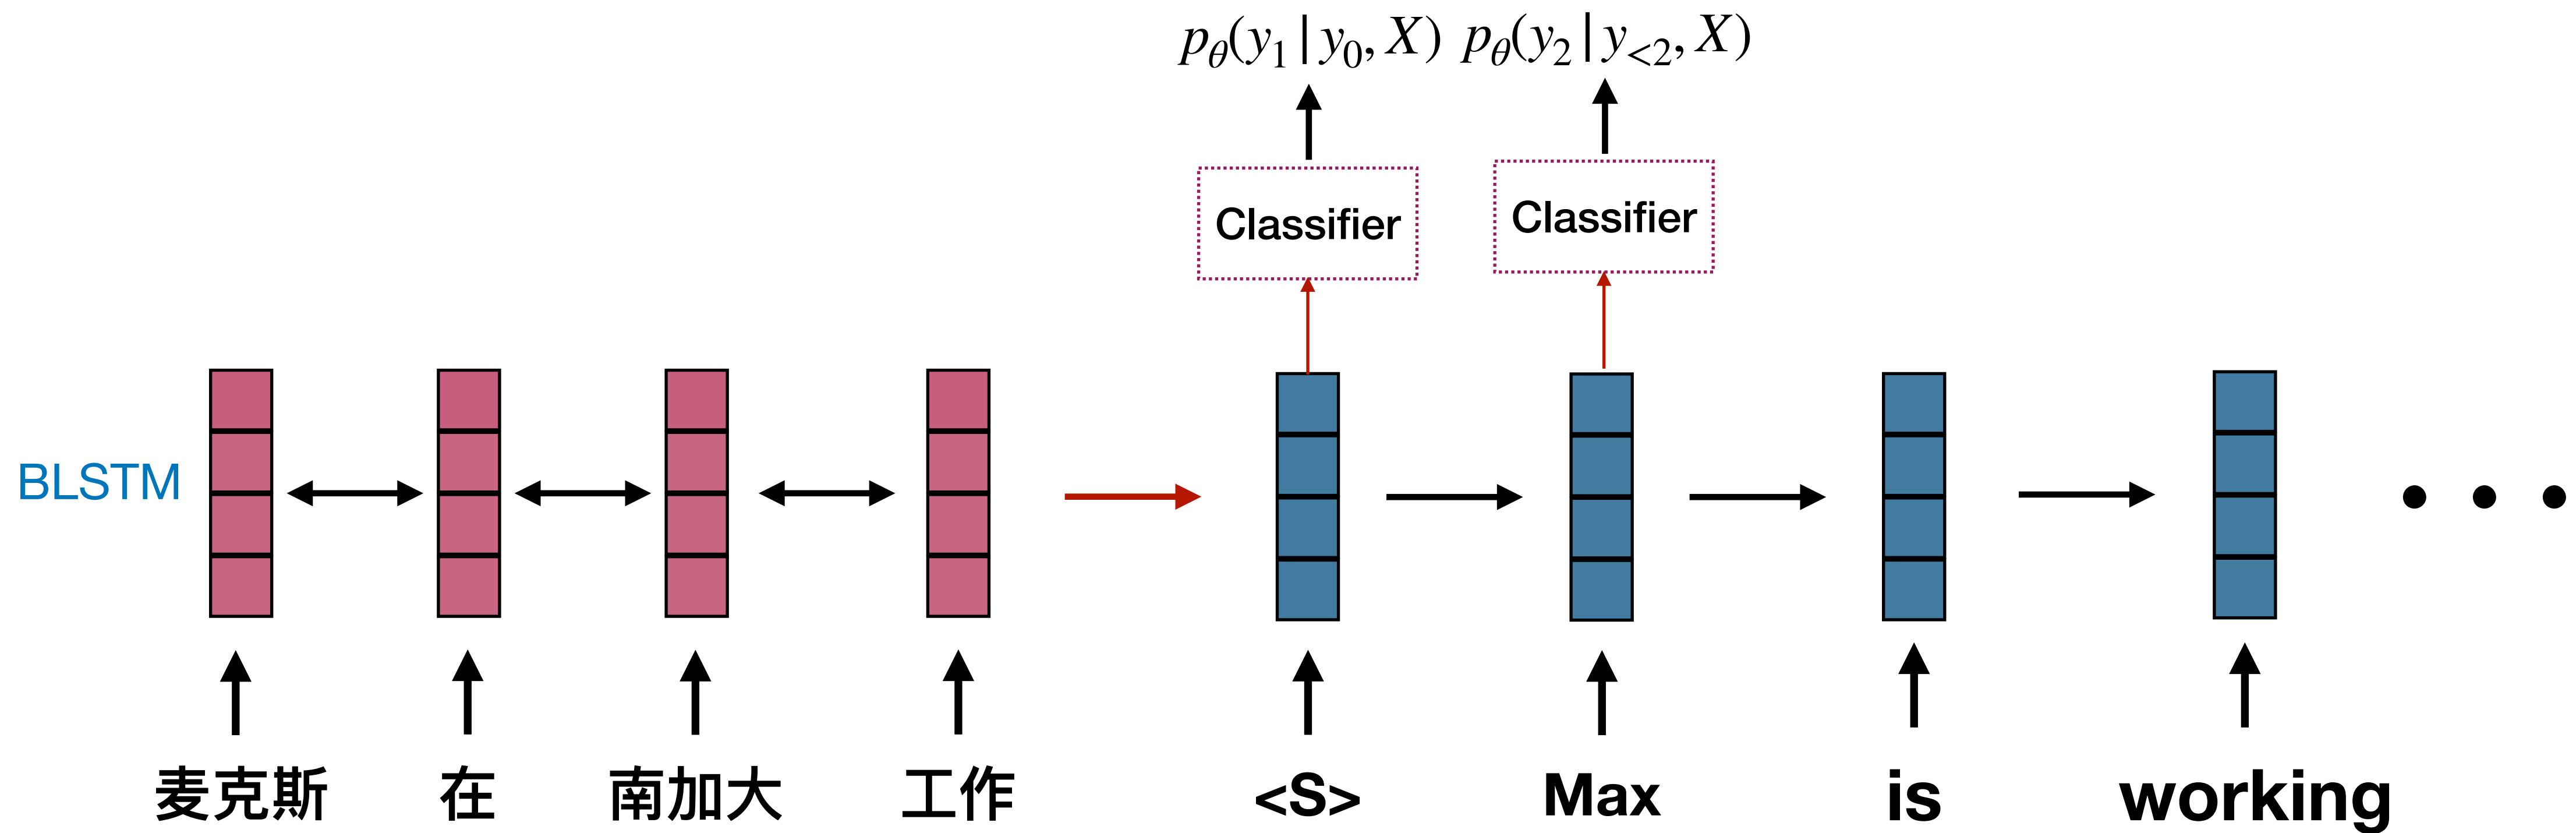
$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X) \quad t = 2$$



# Seq2seq Training

- Model Training:

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$



# Seq2seq Training

- Maximum Likelihood Estimation

$$\max_{\theta} p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$

- Back-propagate gradients through both decoder & encoder
- Need a really big training corpus
  - WMT Russian-English





# Seq2seq Decoding

- Exhaustive Search
  - Requires computing all possible sequences

$$\arg \max_Y p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$

What is the complexity of doing this search, if  $|\mathcal{Y}| = V$  and sequence length  $T$ ?

- (a)  $O(VT)$
- (b)  $O(V^T)$
- (c)  $O(T^V)$

# Seq2seq Decoding

- Greedy Search

- Selects the best current word  $y_t$

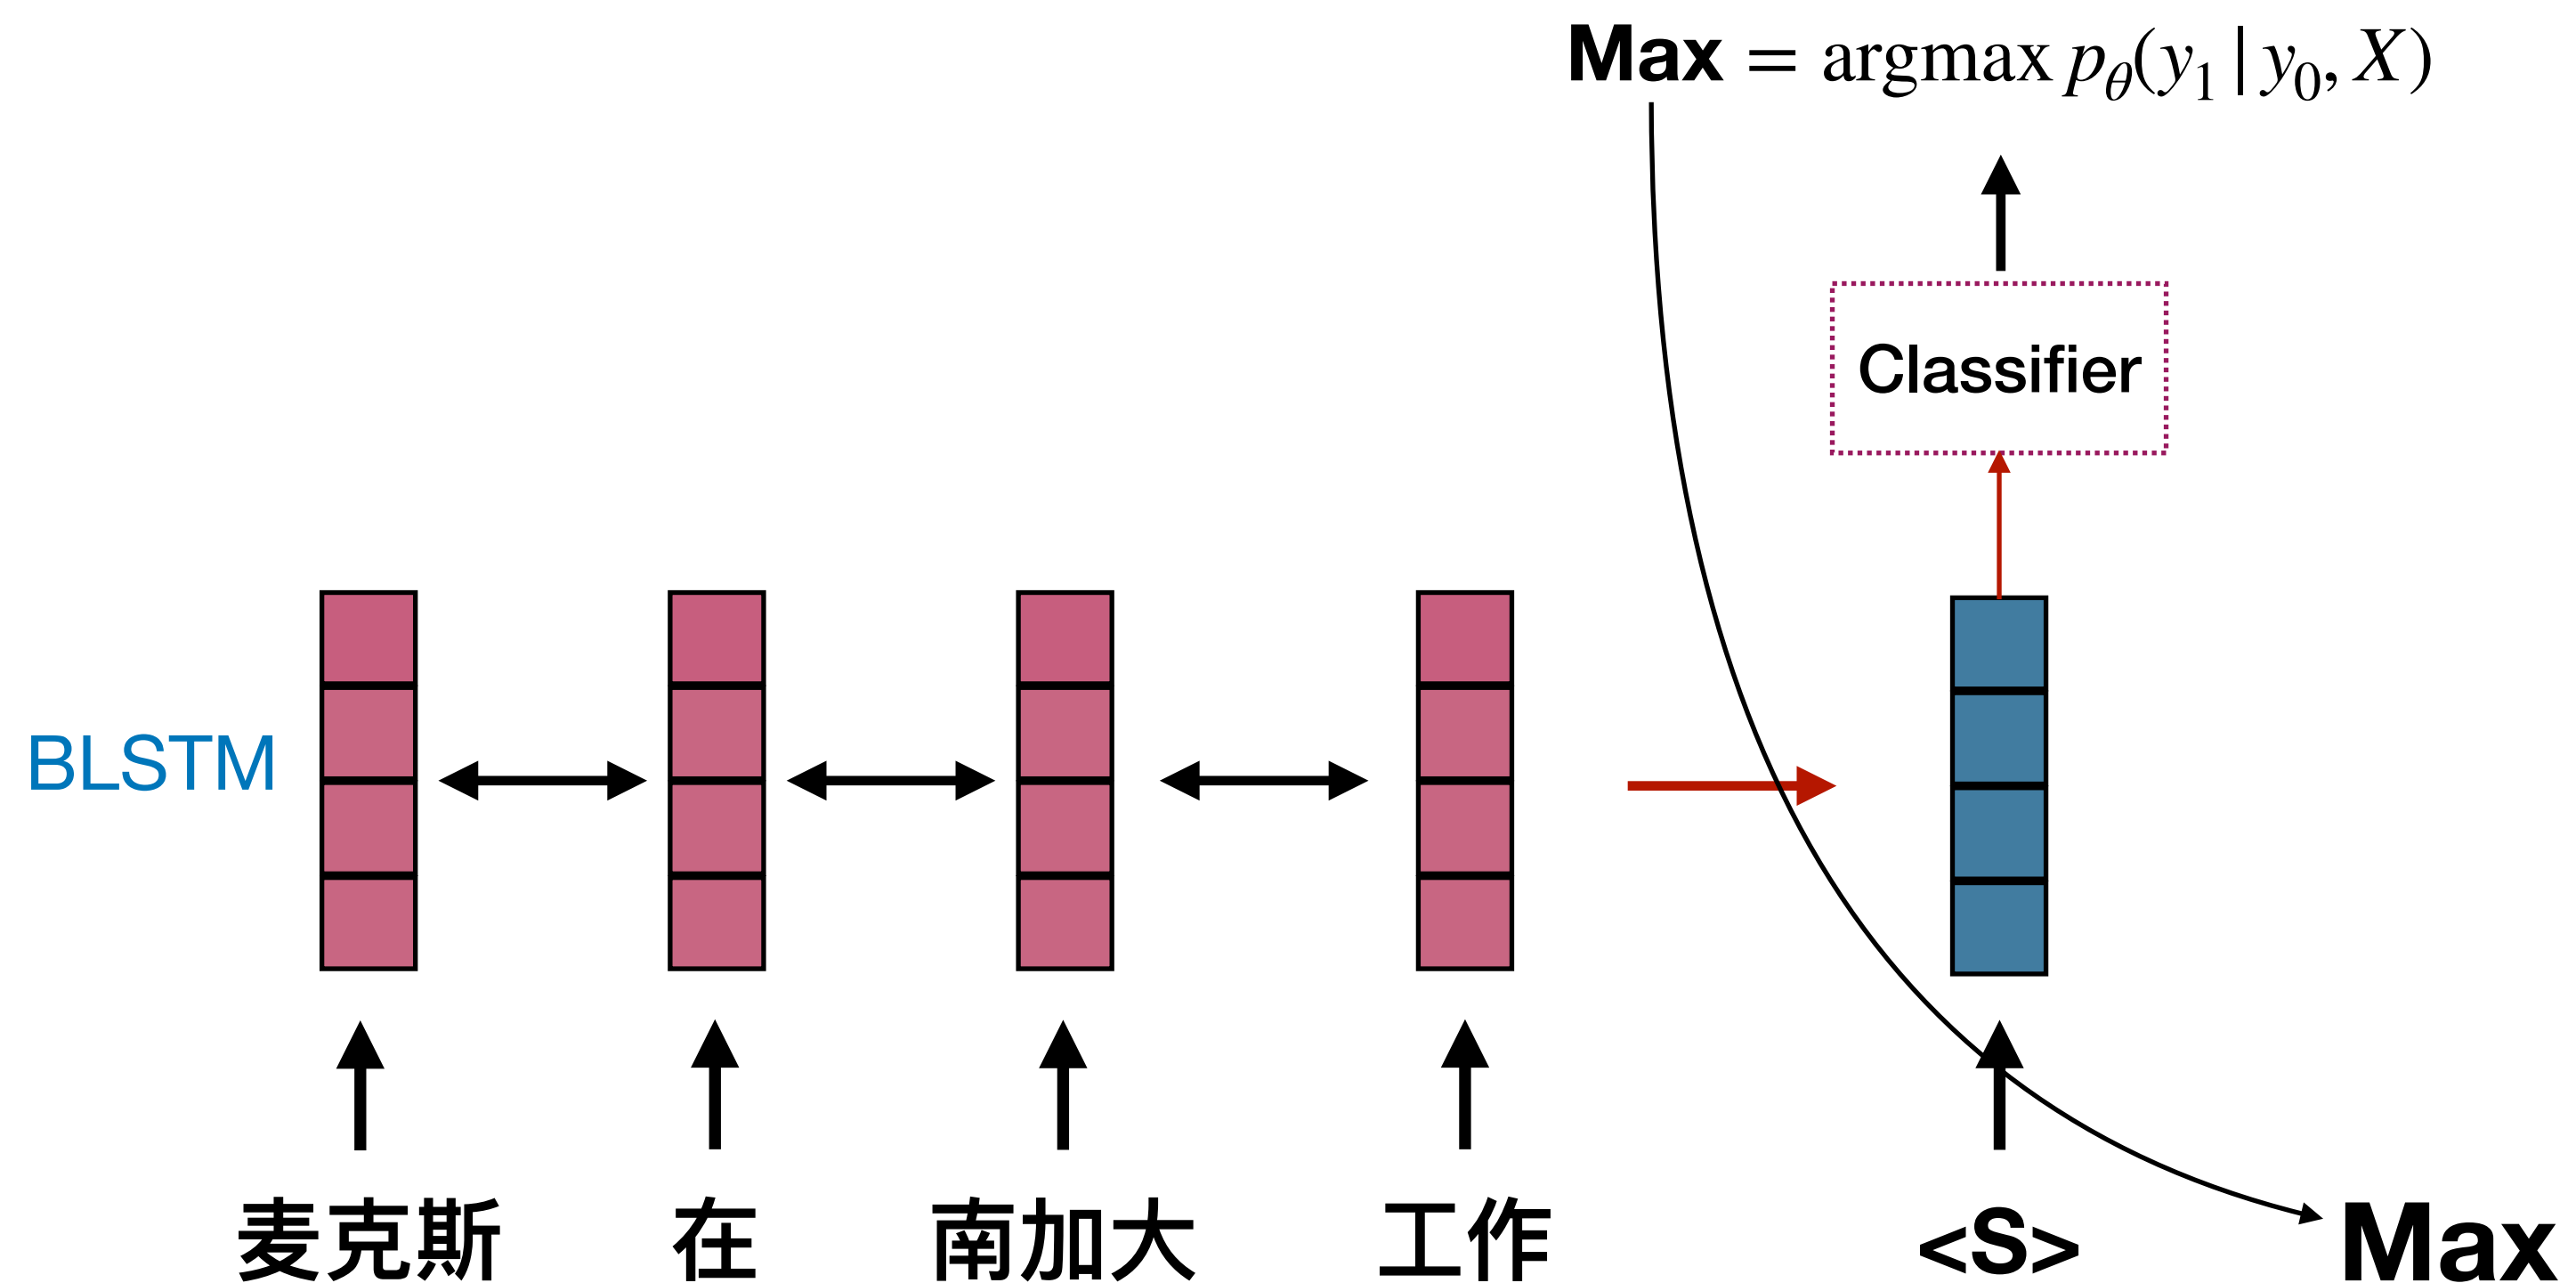
$$\arg \max_Y p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$

$$\approx \mathbf{arg} \max_{y_t} p_{\theta}(y_t | y_{<t}, X), \forall t$$

# Seq2seq Decoding

- Greedy decoding:

$$y_t^* = \arg \max_{y_t} p_{\theta}(y_t | y_{<t}, X), \forall t$$

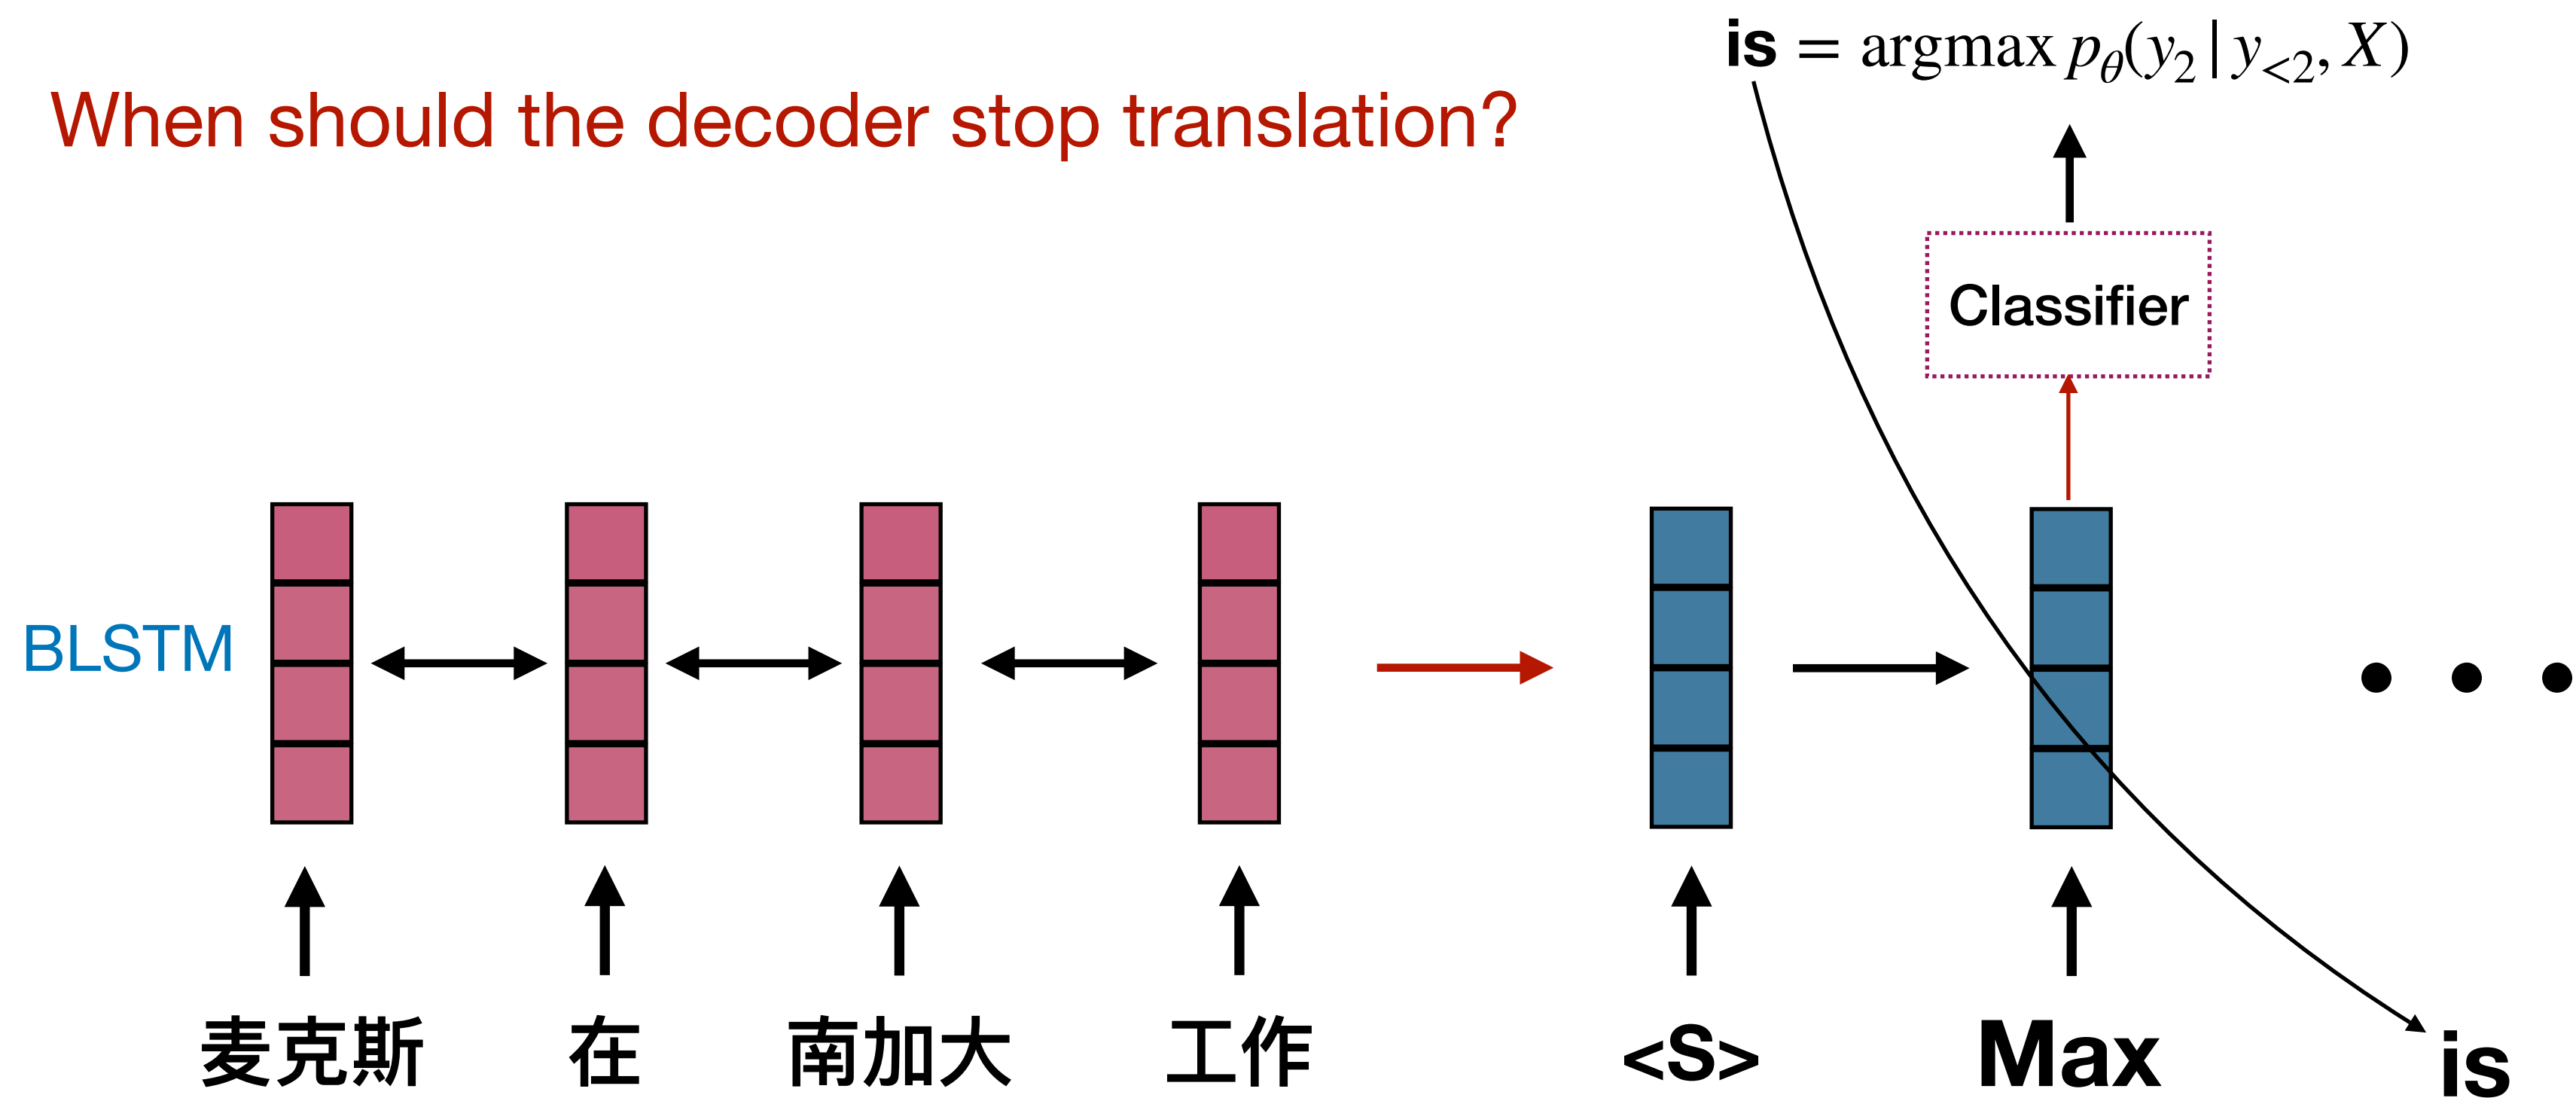


# Seq2seq Decoding

- Greedy decoding:

$$y_t^* = \arg \max_{y_t} p_{\theta}(y_t | y_{<t}, X), \forall t$$

When should the decoder stop translation?



# Special Tokens in Seq2seq

- **<BOS>**: start of the target sentence
- **<EOS>**: end of the target sentence

麦克斯 在 南加大 工作

$x_1$   $x_2$   $x_3$   $x_4$



Max is working at USC

$y_1$   $y_2$   $y_3$   $y_4$   $y_5$

麦克斯 在 南加大 工作

$x_1$   $x_2$   $x_3$   $x_4$



<BOS> Max is working at USC <EOS>


$y_0$   $y_1$   $y_2$   $y_3$   $y_4$   $y_5$   $y_6$

# A Middle Ground: Beam Search

- **Key idea:** at every step, keep track of the **k most probable** partial translations (hypotheses)
- Score of each hypothesis = log probability of sequence so far
- Not guaranteed to be optimal
- More efficient than exhaustive search


# Beam Search Decoding

Beam size  $K = 2$



# Beam Search Decoding


Beam size  $K = 2$





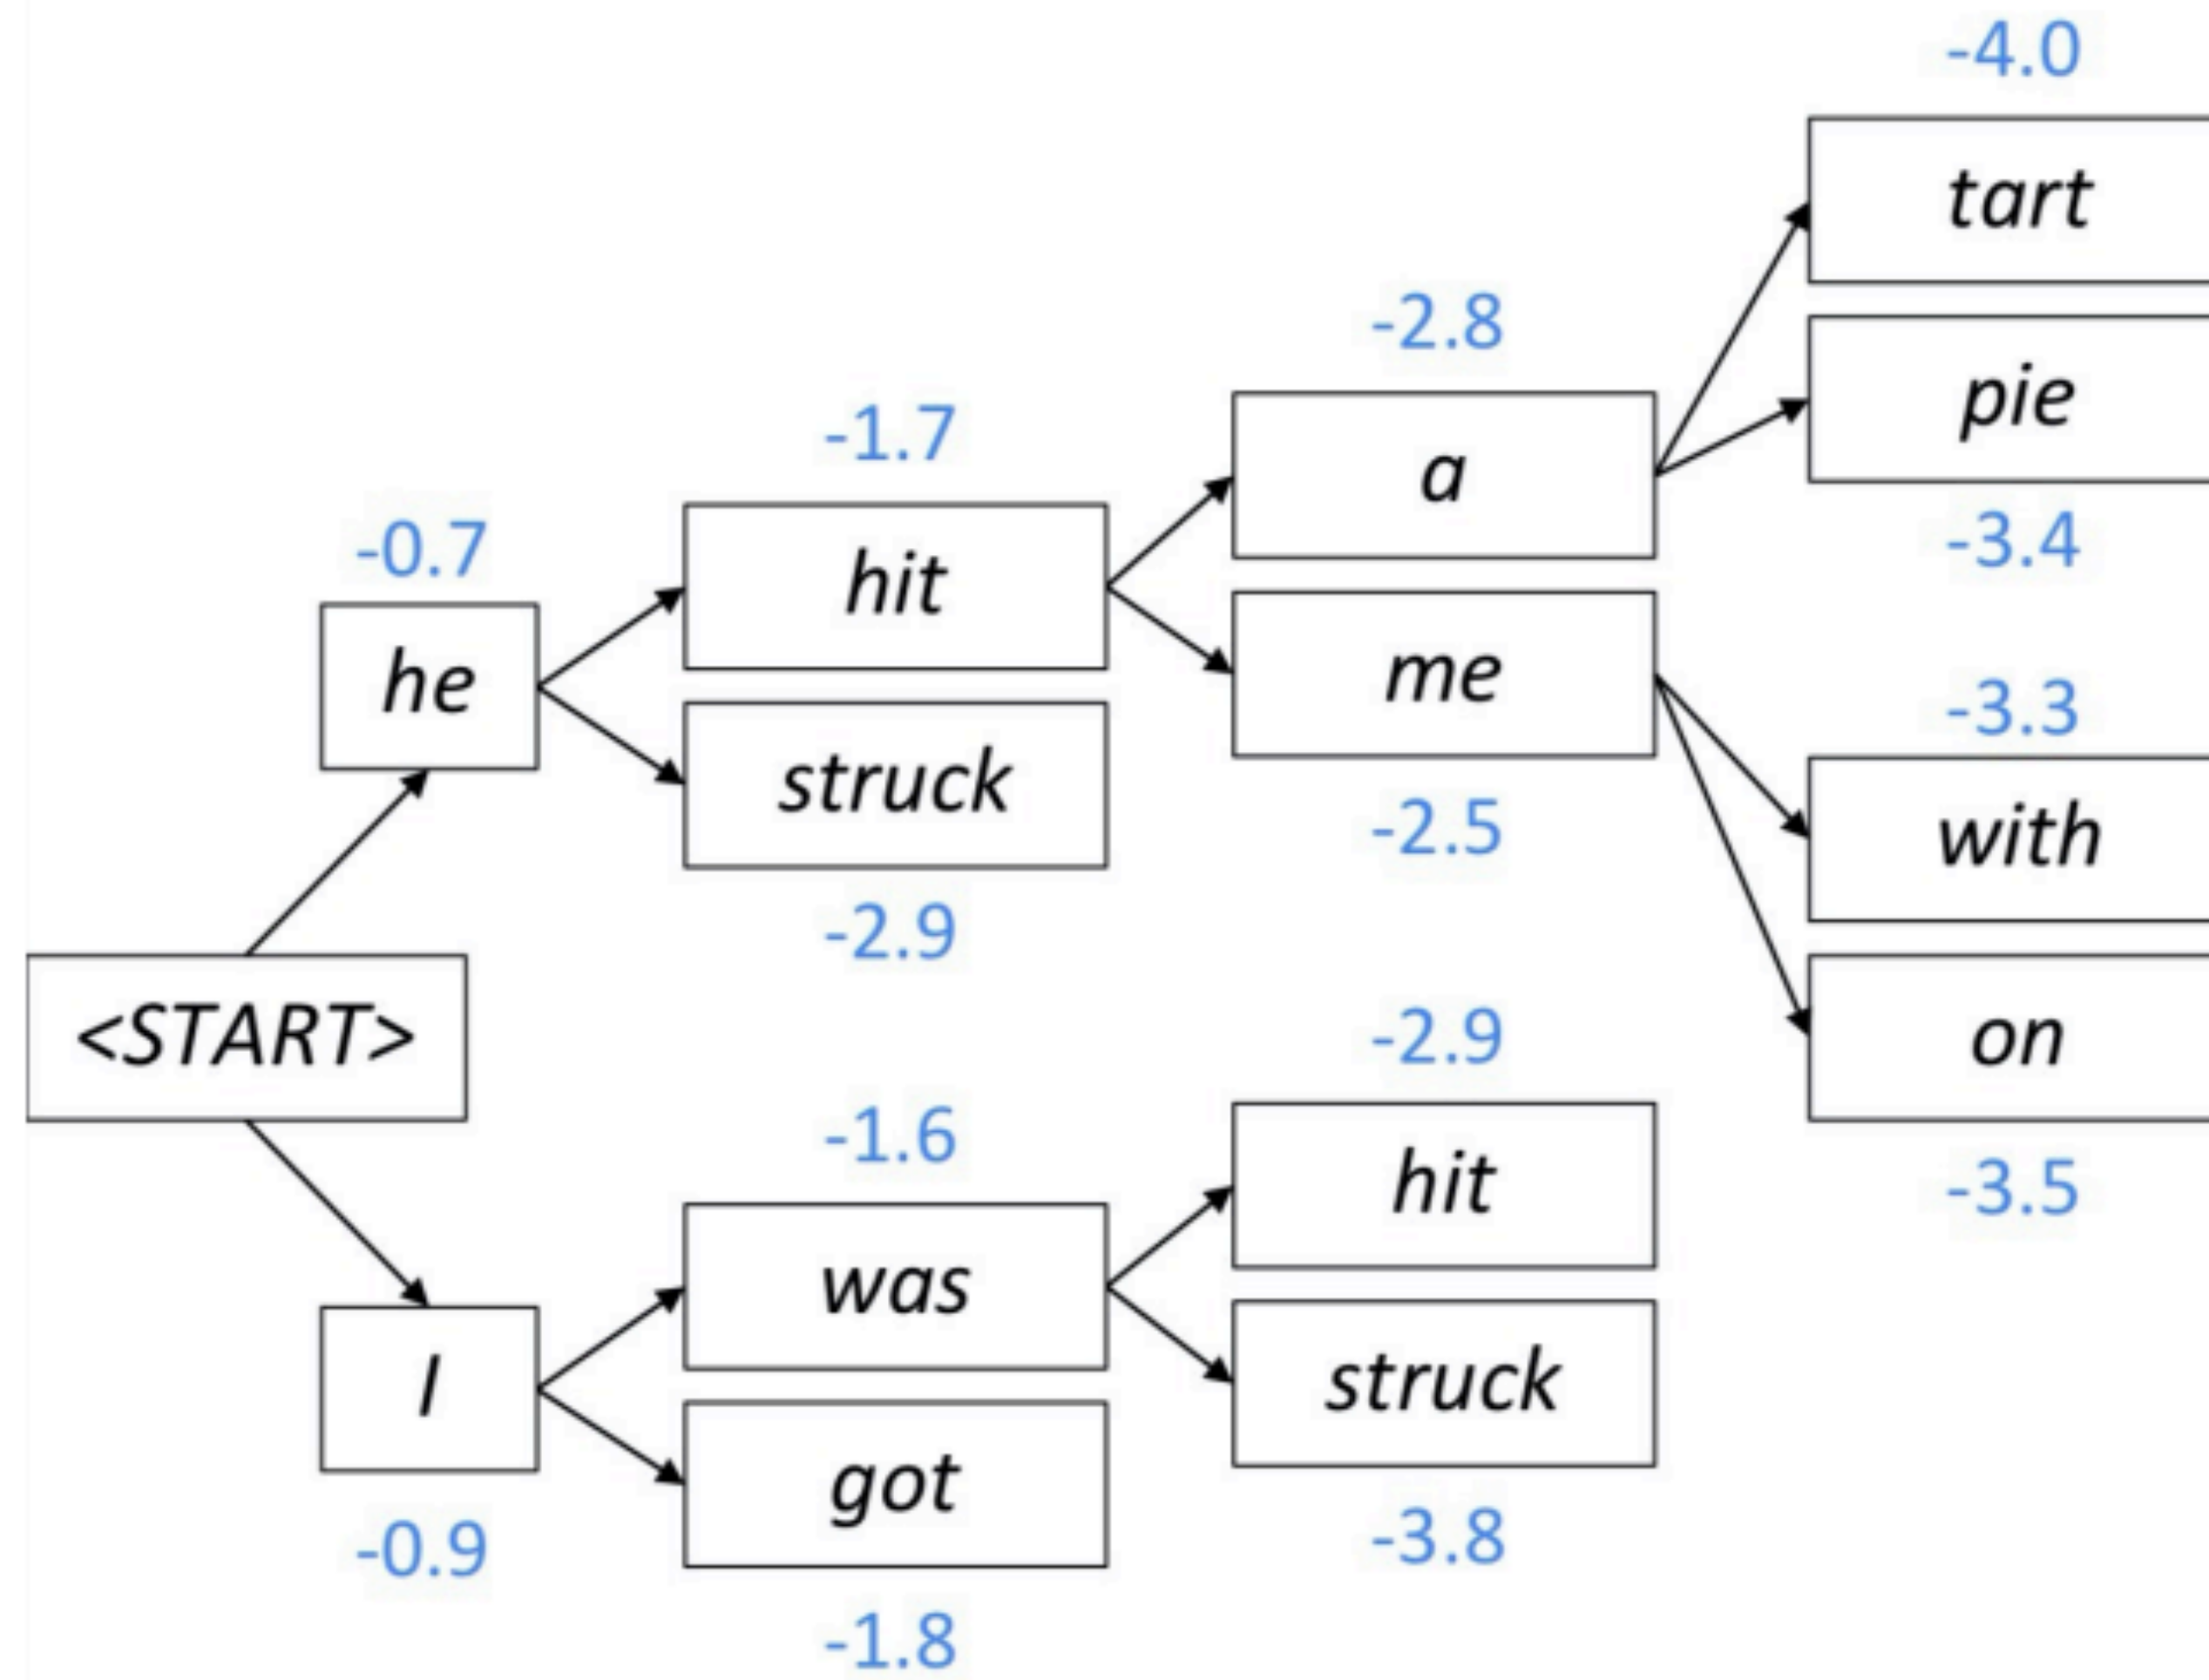
# Beam Search Decoding

Beam size  $K = 2$




# Beam Search Decoding

Beam size  $K = 2$



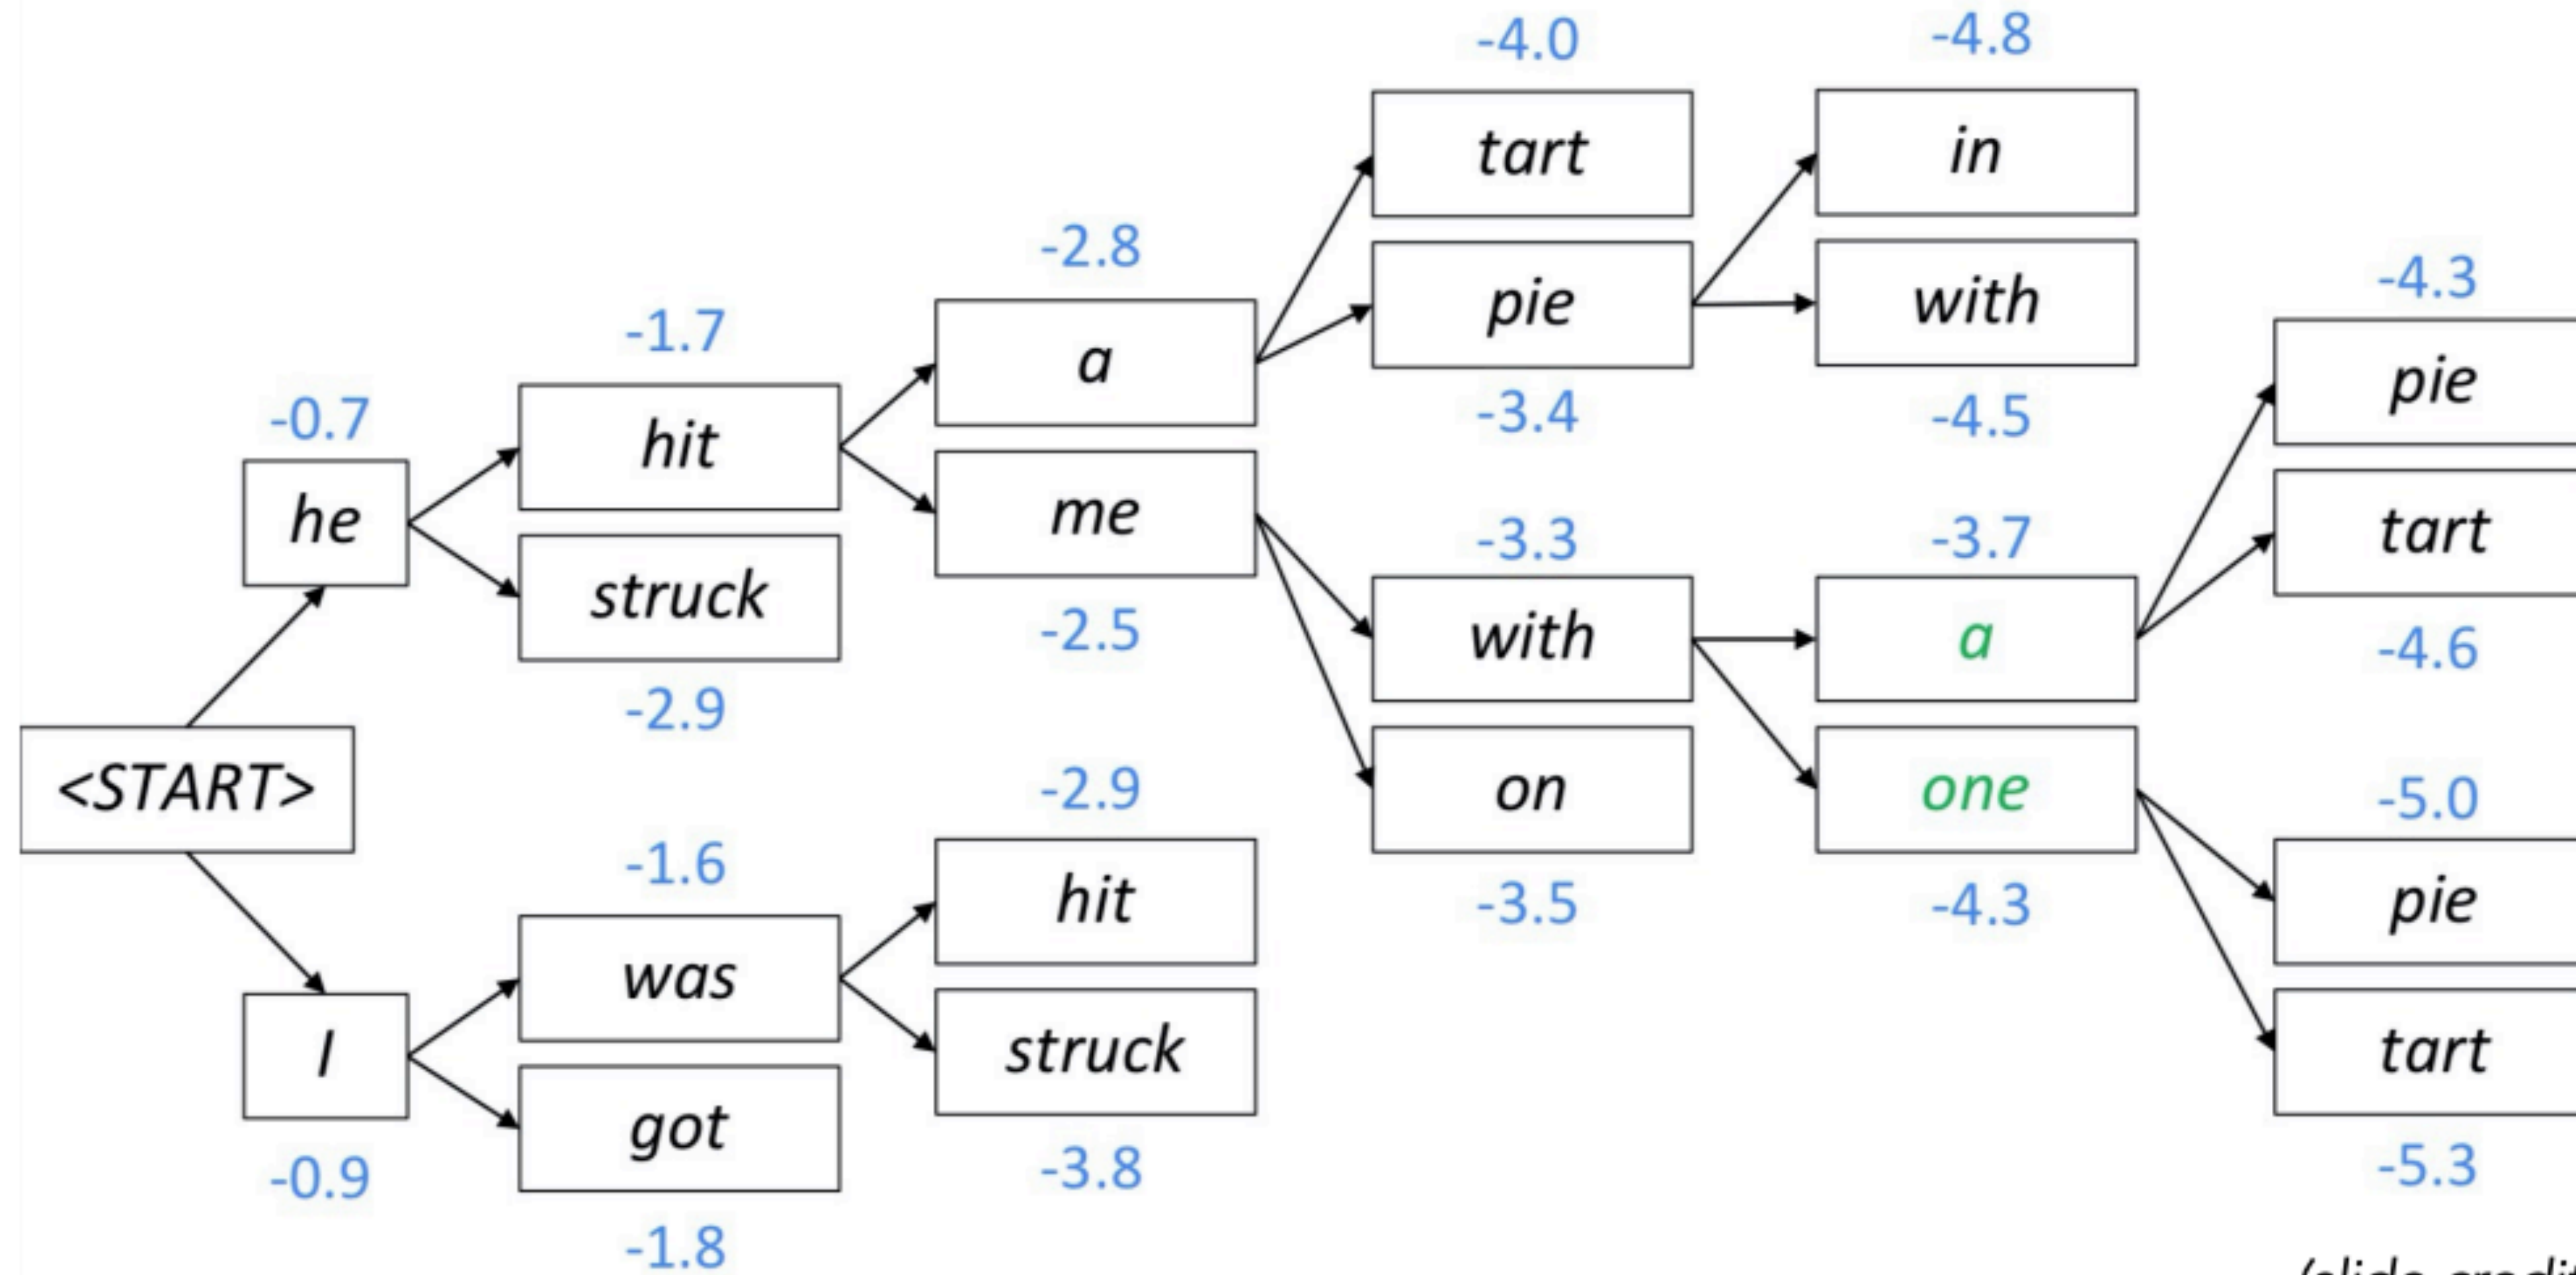
# Beam Search Decoding

Beam size  $K = 2$



# Beam Search Decoding

Beam size  $K = 2$

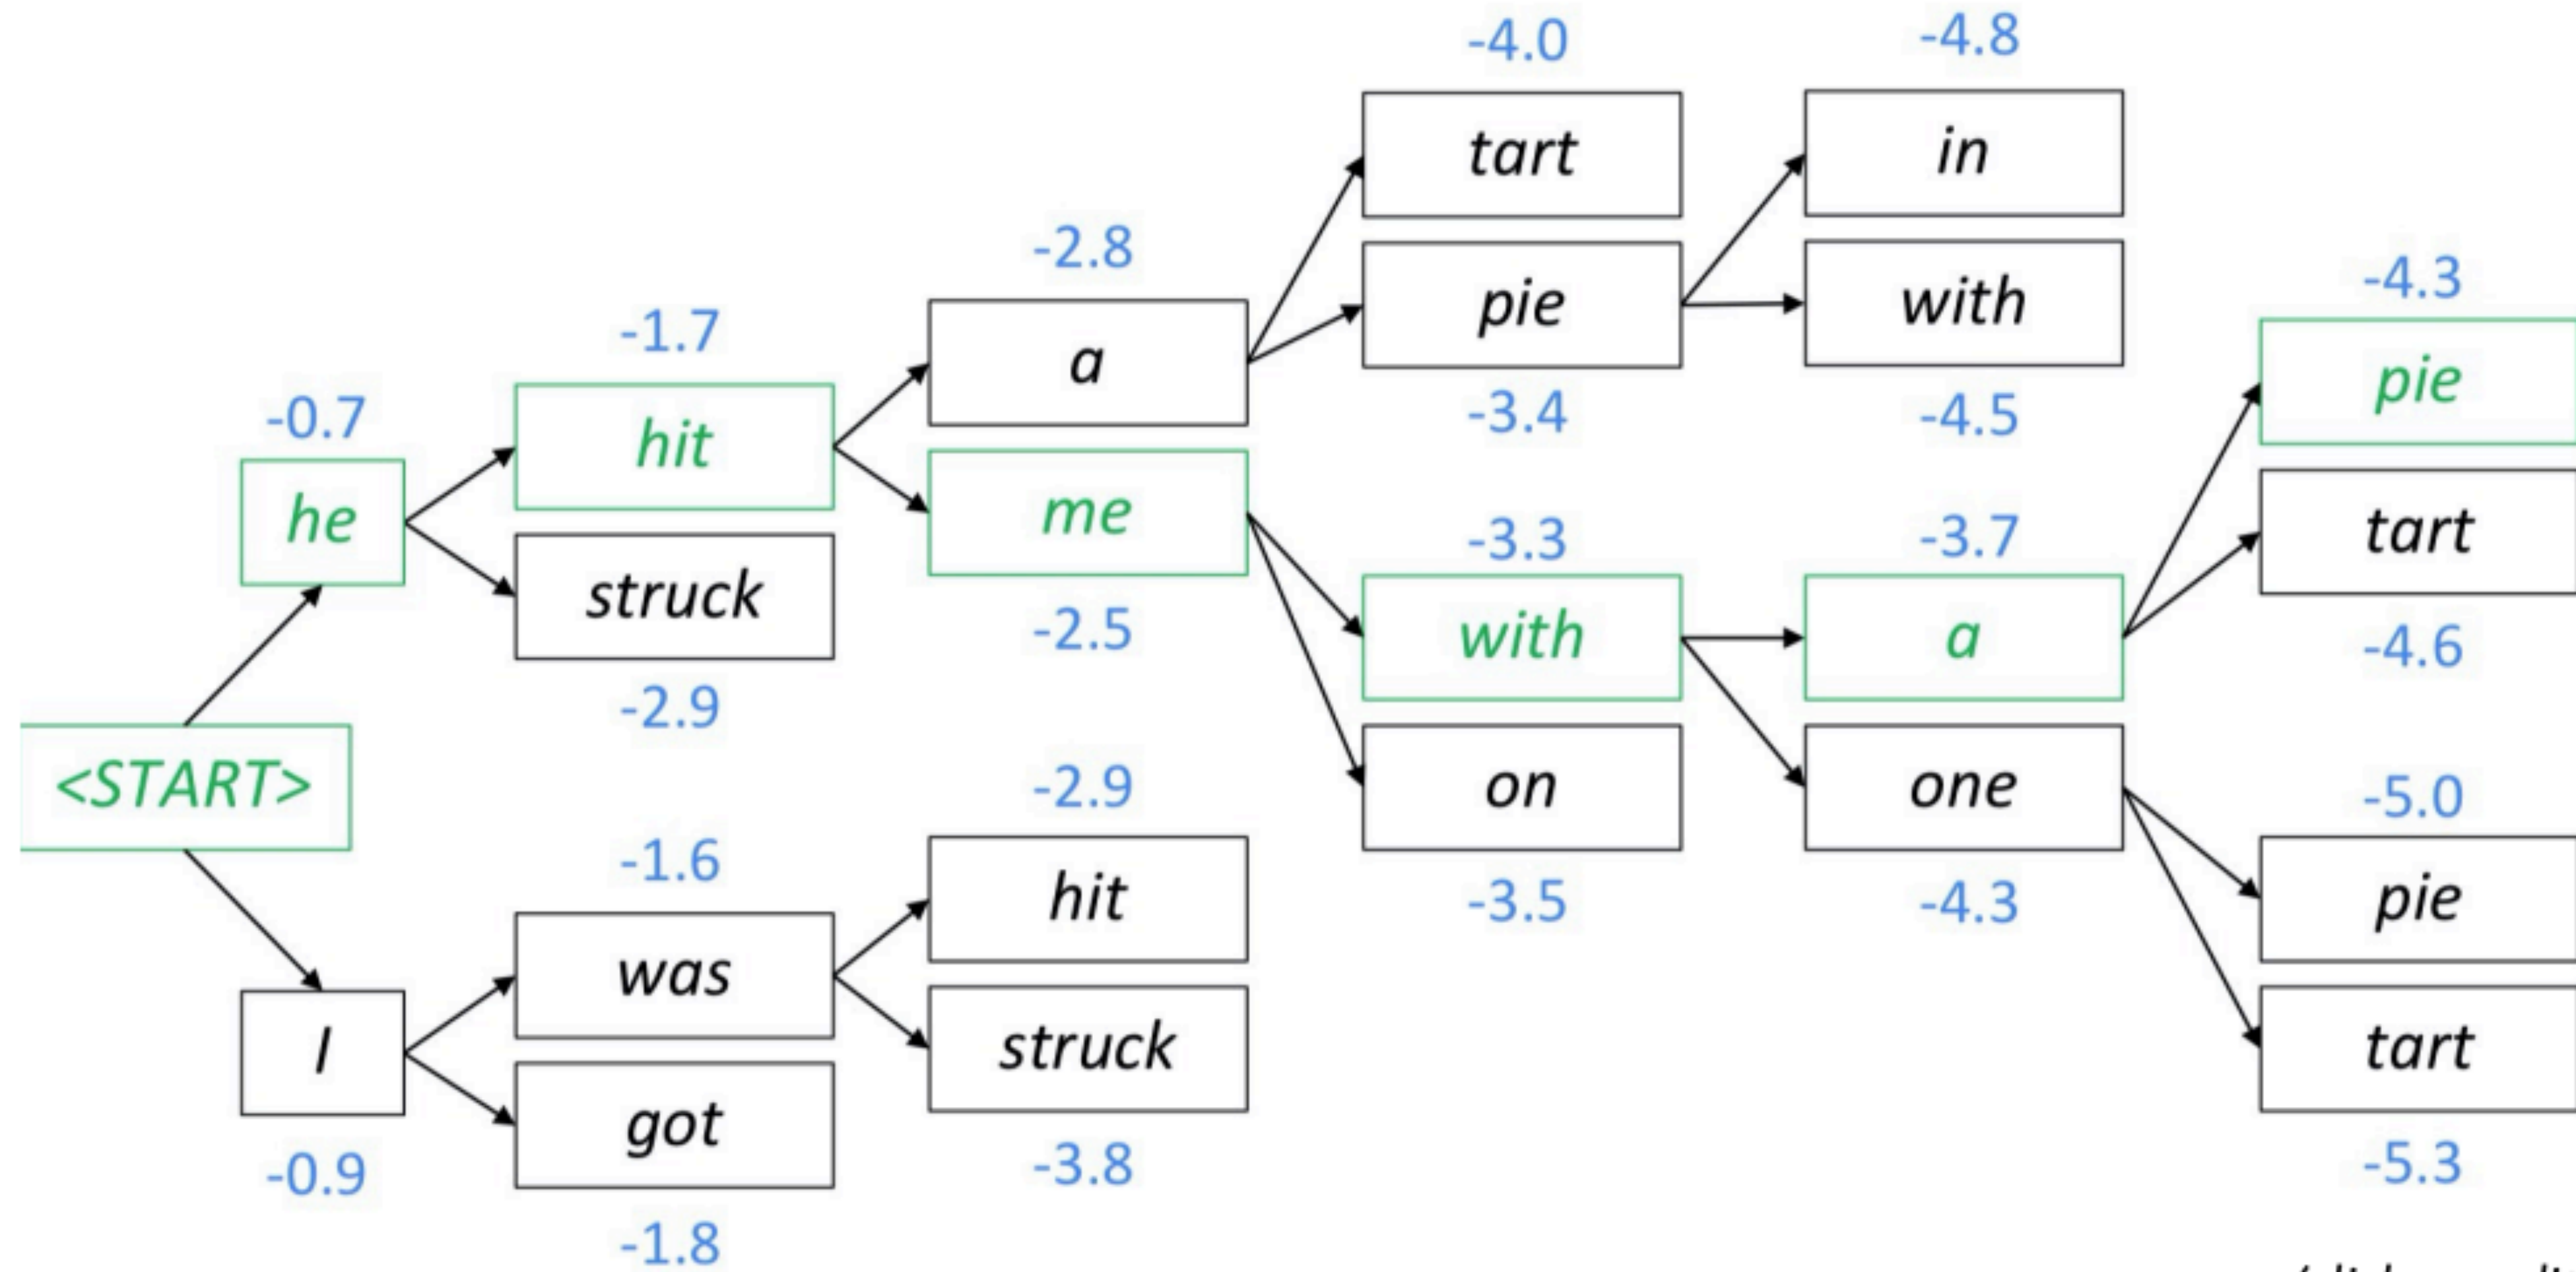


(slide credit: Abigail See)




# Backtrack

Beam size  $K = 2$




(slide credit: Abigail See)

# How Seq2seq changed the MT Landscape




# MT Progress



# Issues with Vanilla Encoder-Decoder Architecture


- A single encoding vector needs to capture **all the information** about source sentence






# Issues with Vanilla Encoder-Decoder Architecture

- A single encoding vector needs to capture **all the information** about source sentence
- Longer sequences can lead to **vanishing gradients**



# Attention Mechanism


- **Key idea:** At each time step, use all parts of source sentence



# Attention Mechanism



$$= \text{attn}([c_1, c_2, \dots, c_n], h_t)$$




# Attention Mechanism



$$= \text{attn}([c_1, c_2, \dots, c_n], h_t)$$

$$e_j^t = \text{sim}(c_j, h_t), \forall j \in \{1, \dots, n\}$$

Attention scores



# Attention Mechanism




$$= \text{attn}([c_1, c_2, \dots, c_n], h_t)$$

$$a^t = \text{softmax}(e^t) \in (0,1)^n$$

Attention distribution

$$e_j^t = \text{sim}(c_j, h_t), \forall j \in \{1, \dots, n\}$$

Attention scores



# Attention Mechanism



$$= \text{attn}([c_1, c_2, \dots, c_n], h_t) = \sum_{j=1}^n a_j^t c_j \in \mathbb{R}^d$$


Attention output

$$a^t = \text{softmax}(e^t) \in (0,1)^n$$

Attention distribution

$$e_j^t = \text{sim}(c_j, h_t), \forall j \in \{1, \dots, n\}$$

Attention scores



# Softmax Function

$$e^t = [e_1^t, e_2^t, \dots, e_n^t]$$

$$\text{softmax}(e^t) = \left[ \frac{\exp(e_1^t)}{\sum_{j=1}^n \exp(e_j^t)}, \frac{\exp(e_2^t)}{\sum_{j=1}^n \exp(e_j^t)}, \dots, \frac{\exp(e_n^t)}{\sum_{j=1}^n \exp(e_j^t)} \right]$$

# Attention Mechanism



$$= \text{attn}([c_1, c_2, \dots, c_n], h_t) = \sum_{j=1}^n a_j^t c_j \in \mathbb{R}^d$$


Attention output

$$a^t = \text{softmax}(e^t) \in (0,1)^n$$

Attention distribution

$$e_j^t = \text{sim}(c_j, h_t), \forall j \in \{1, \dots, n\}$$

Attention scores





# Types of Attention

- **Dot-product attention** (assumes equal dimensions for  $c$  and  $h$ )

$$\text{sim}(c_j, h_t) = c_j^T h_t$$

- **Multiplicative attention**


$$\text{sim}(c_j, h_t) = c_j^T W h_t, \text{ where } W \text{ is learnable weight matrix}$$

- **Additive attention**

$$\text{sim}(c_j, h_t) = v^T \tanh(W_c c_j + W_h h_t)$$

where  $W_c$  and  $W_h$  are learnable weight matrices and  $v$  is a learnable weight vector

# Visualizing Attention



Highly correlated with alignment

# Attention Improves Translation Performance

System	Ppl	BLEU
Winning WMT'14 system – <i>phrase-based</i> + <i>large LM</i> (Buck et al., 2014)		20.7
<i>Existing NMT systems</i>		
RNNsearch (Jean et al., 2015)		16.5
RNNsearch + unk replace (Jean et al., 2015)		19.0
RNNsearch + unk replace + large vocab + <i>ensemble</i> 8 models (Jean et al., 2015)		<b>21.6</b>
<i>Our NMT systems</i>		
Base	10.6	11.3
Base + reverse	9.9	12.6 (+1.3)
Base + reverse + dropout	8.1	14.0 (+1.4)
Base + reverse + dropout + global attention ( <i>location</i> )	7.3	16.8 (+2.8)
Base + reverse + dropout + global attention ( <i>location</i> ) + feed input	6.4	18.1 (+1.3)
Base + reverse + dropout + local-p attention ( <i>general</i> ) + feed input	5.9	19.0 (+0.9)
Base + reverse + dropout + local-p attention ( <i>general</i> ) + feed input + unk replace		20.9 (+1.9)
<i>Ensemble</i> 8 models + unk replace		<b>23.0 (+2.1)</b>

(Luong et al., 2015)




# Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation

Table 10: Mean of side-by-side scores on production data

	PBMT	GNMT	Human	Relative Improvement
English → Spanish	4.885	5.428	5.504	87%
English → French	4.932	5.295	5.496	64%
English → Chinese	4.035	4.594	4.987	58%
Spanish → English	4.872	5.187	5.372	63%
French → English	5.046	5.343	5.404	83%
Chinese → English	3.694	4.263	4.636	60%

# MT Progress



# Reading Materials

- **Reading Materials**
  - Sequence to Sequence Learning with Neural Networks
  - Neural Machine Translation by Jointly Learning to Align and Translate

