


CSCI 544: Applied Natural Language Processing

Transformer-I

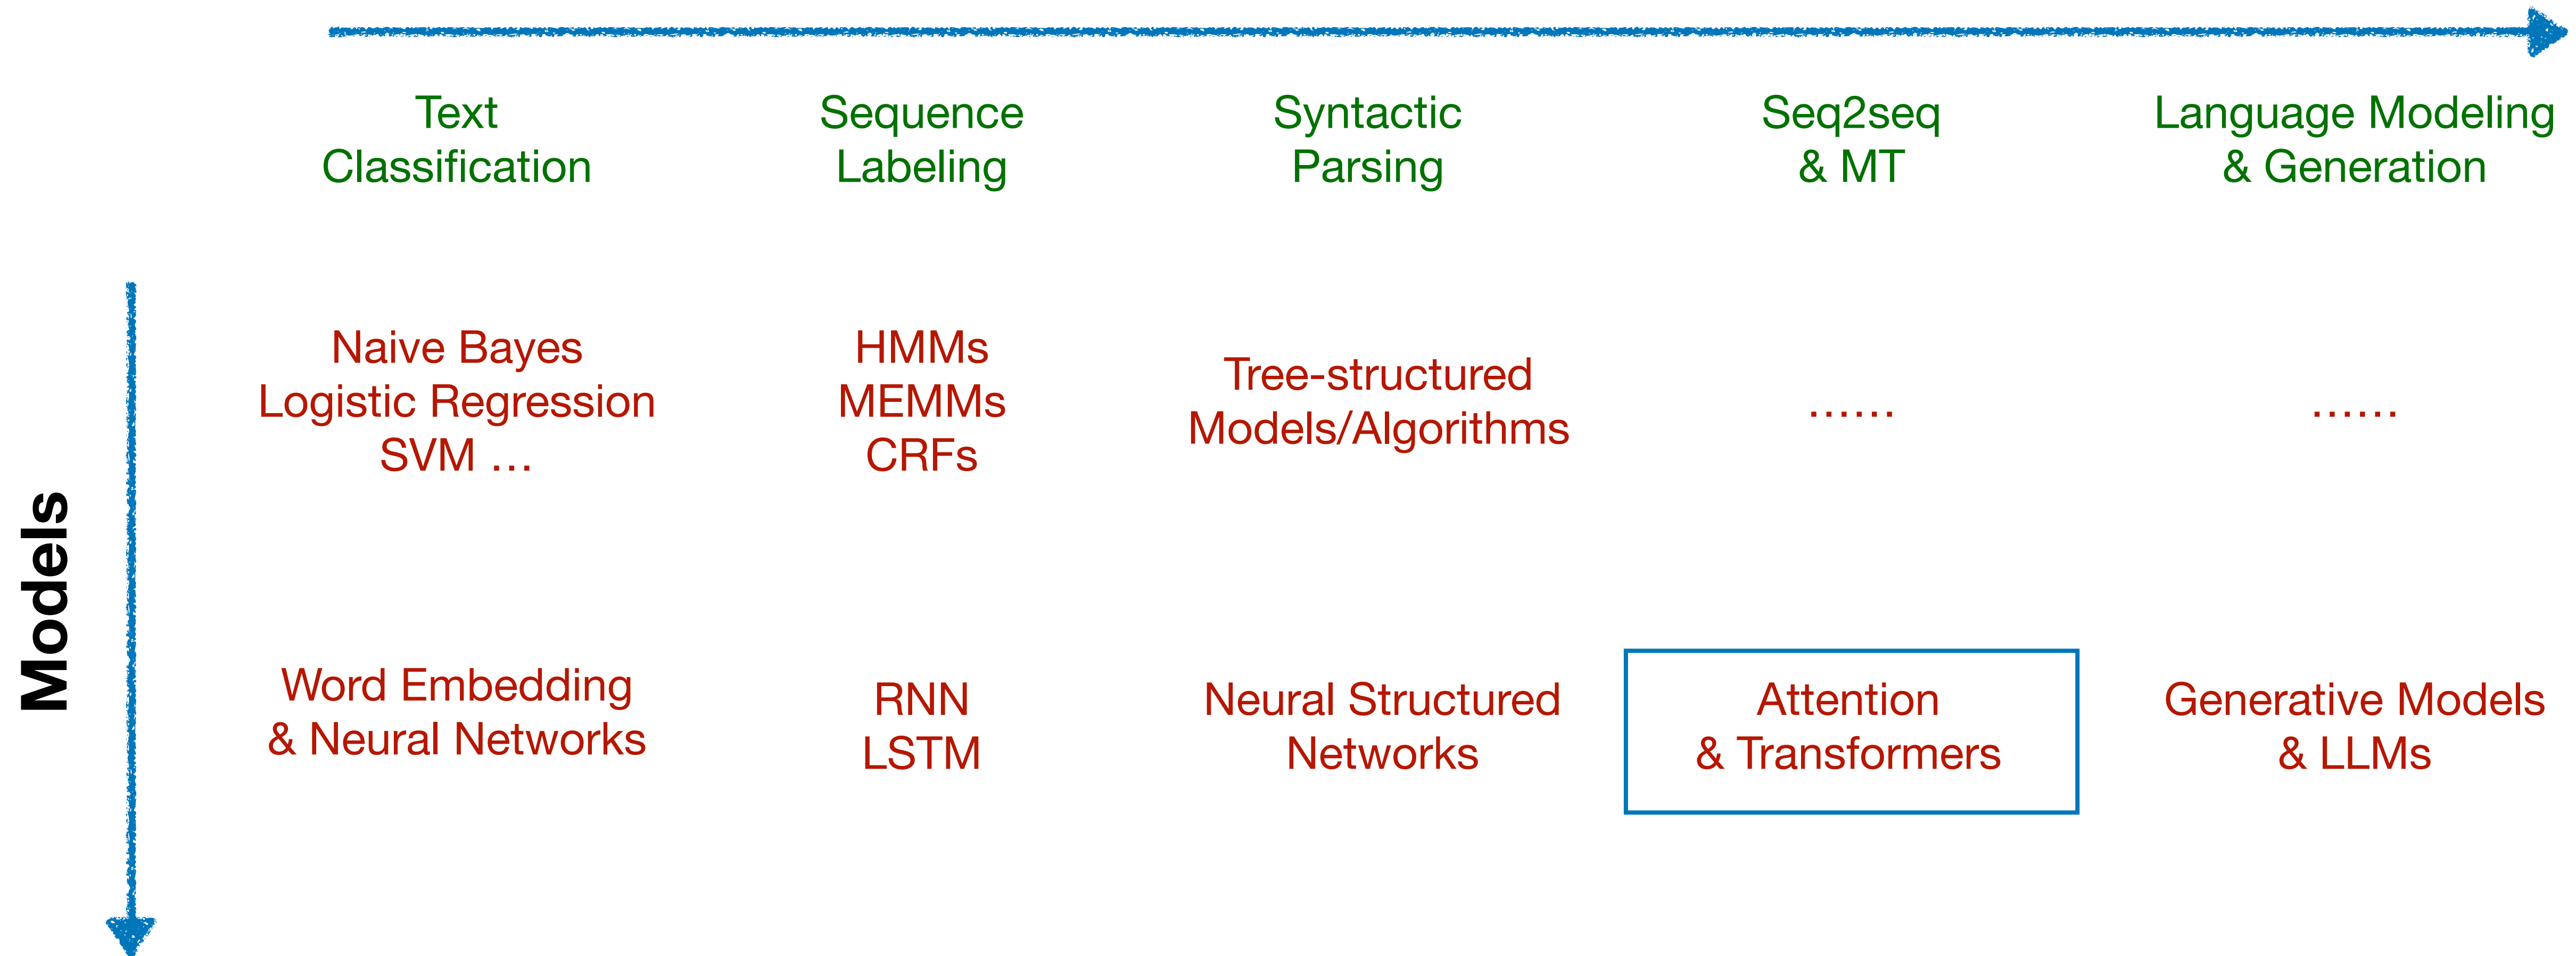
Xuezhe Ma (Max)



USC University of
Southern California

Course Organization

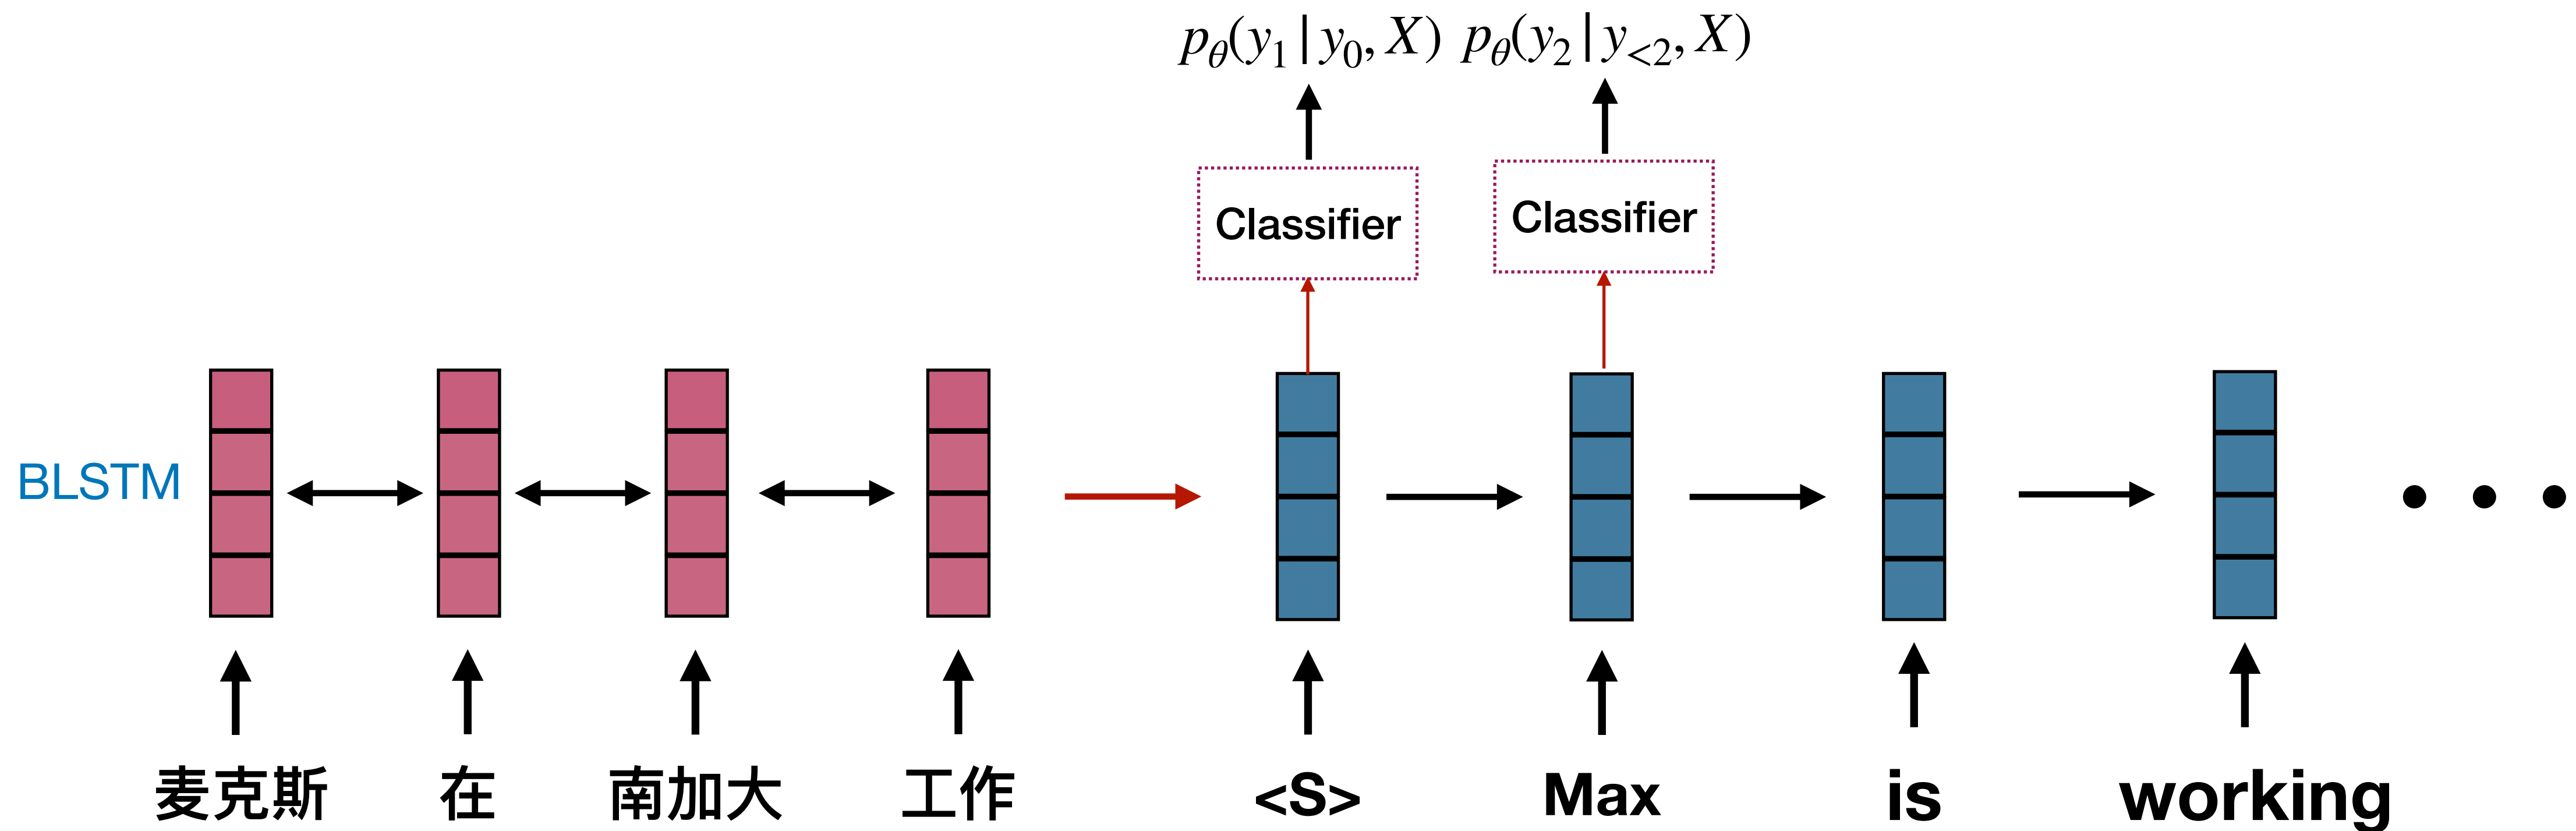
NLP Tasks



Recap: Encoder-Decoder Architecture

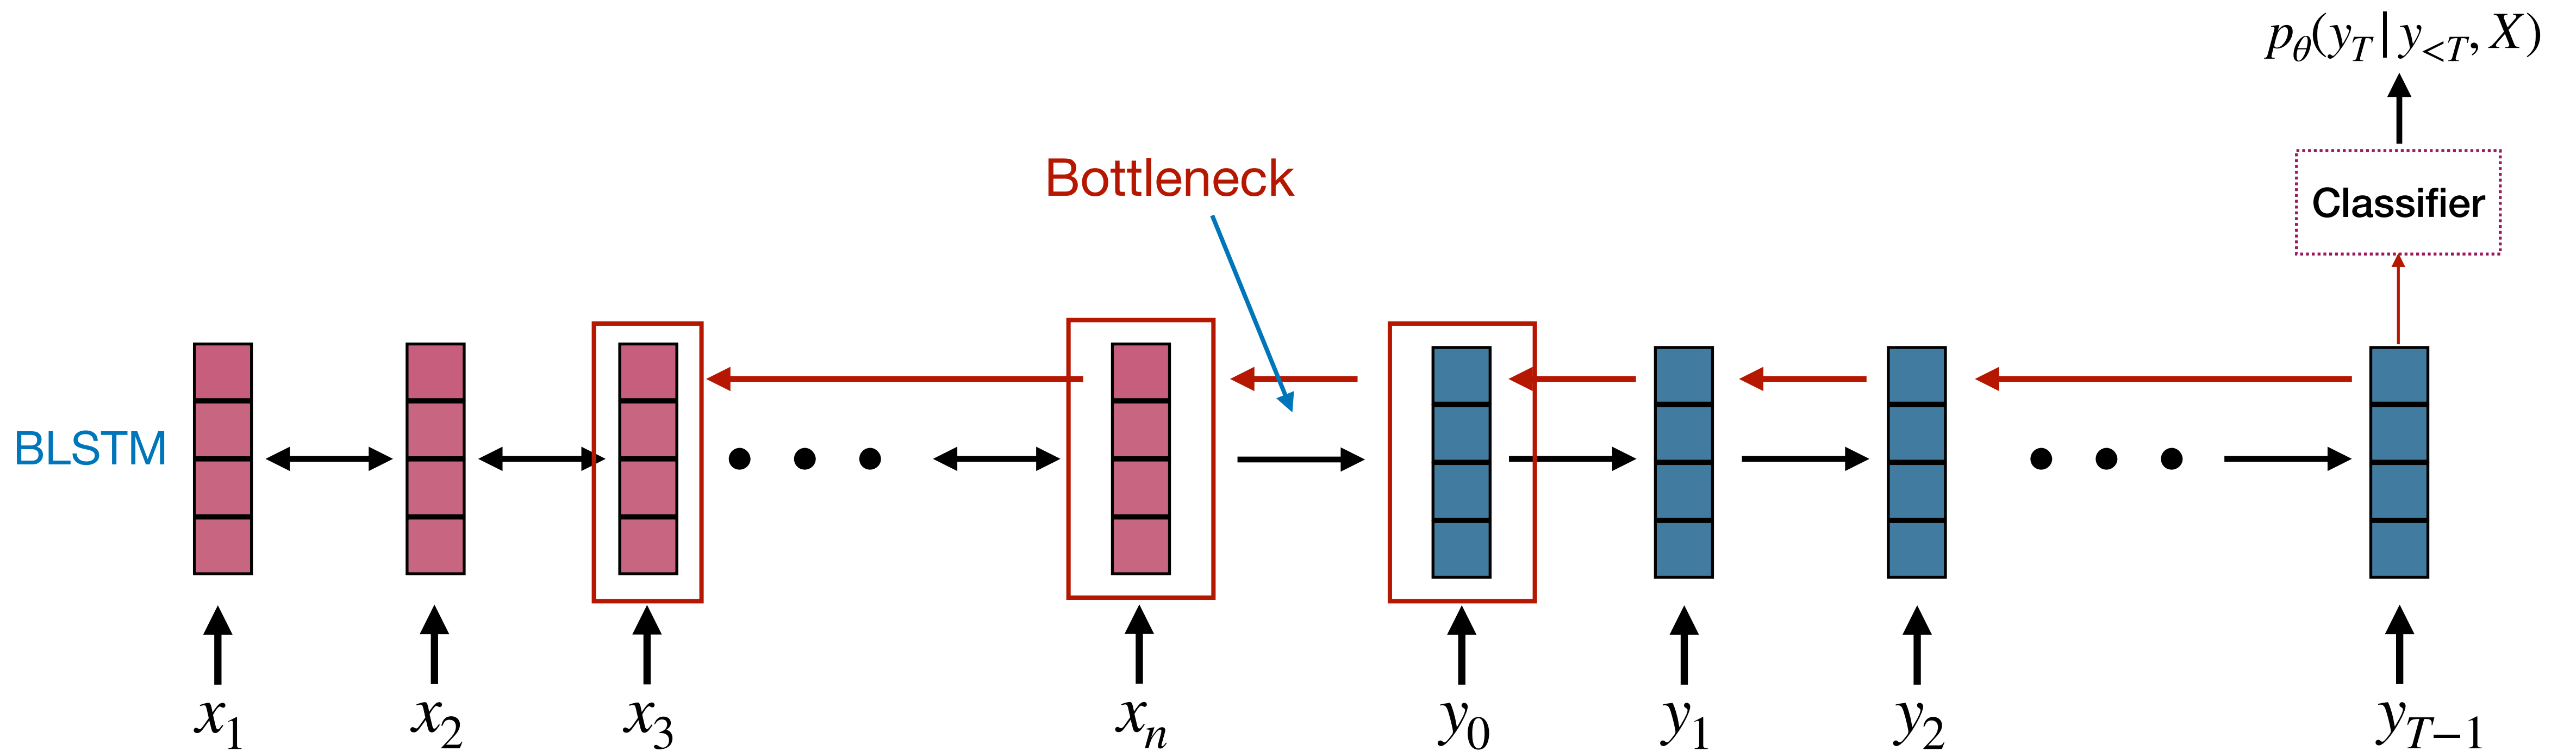
- **Two Components:**

- **Encoder:** Convert input sequence into a sequence of vectors
- **Decoder:** Convert decoding into a sequence in the output space

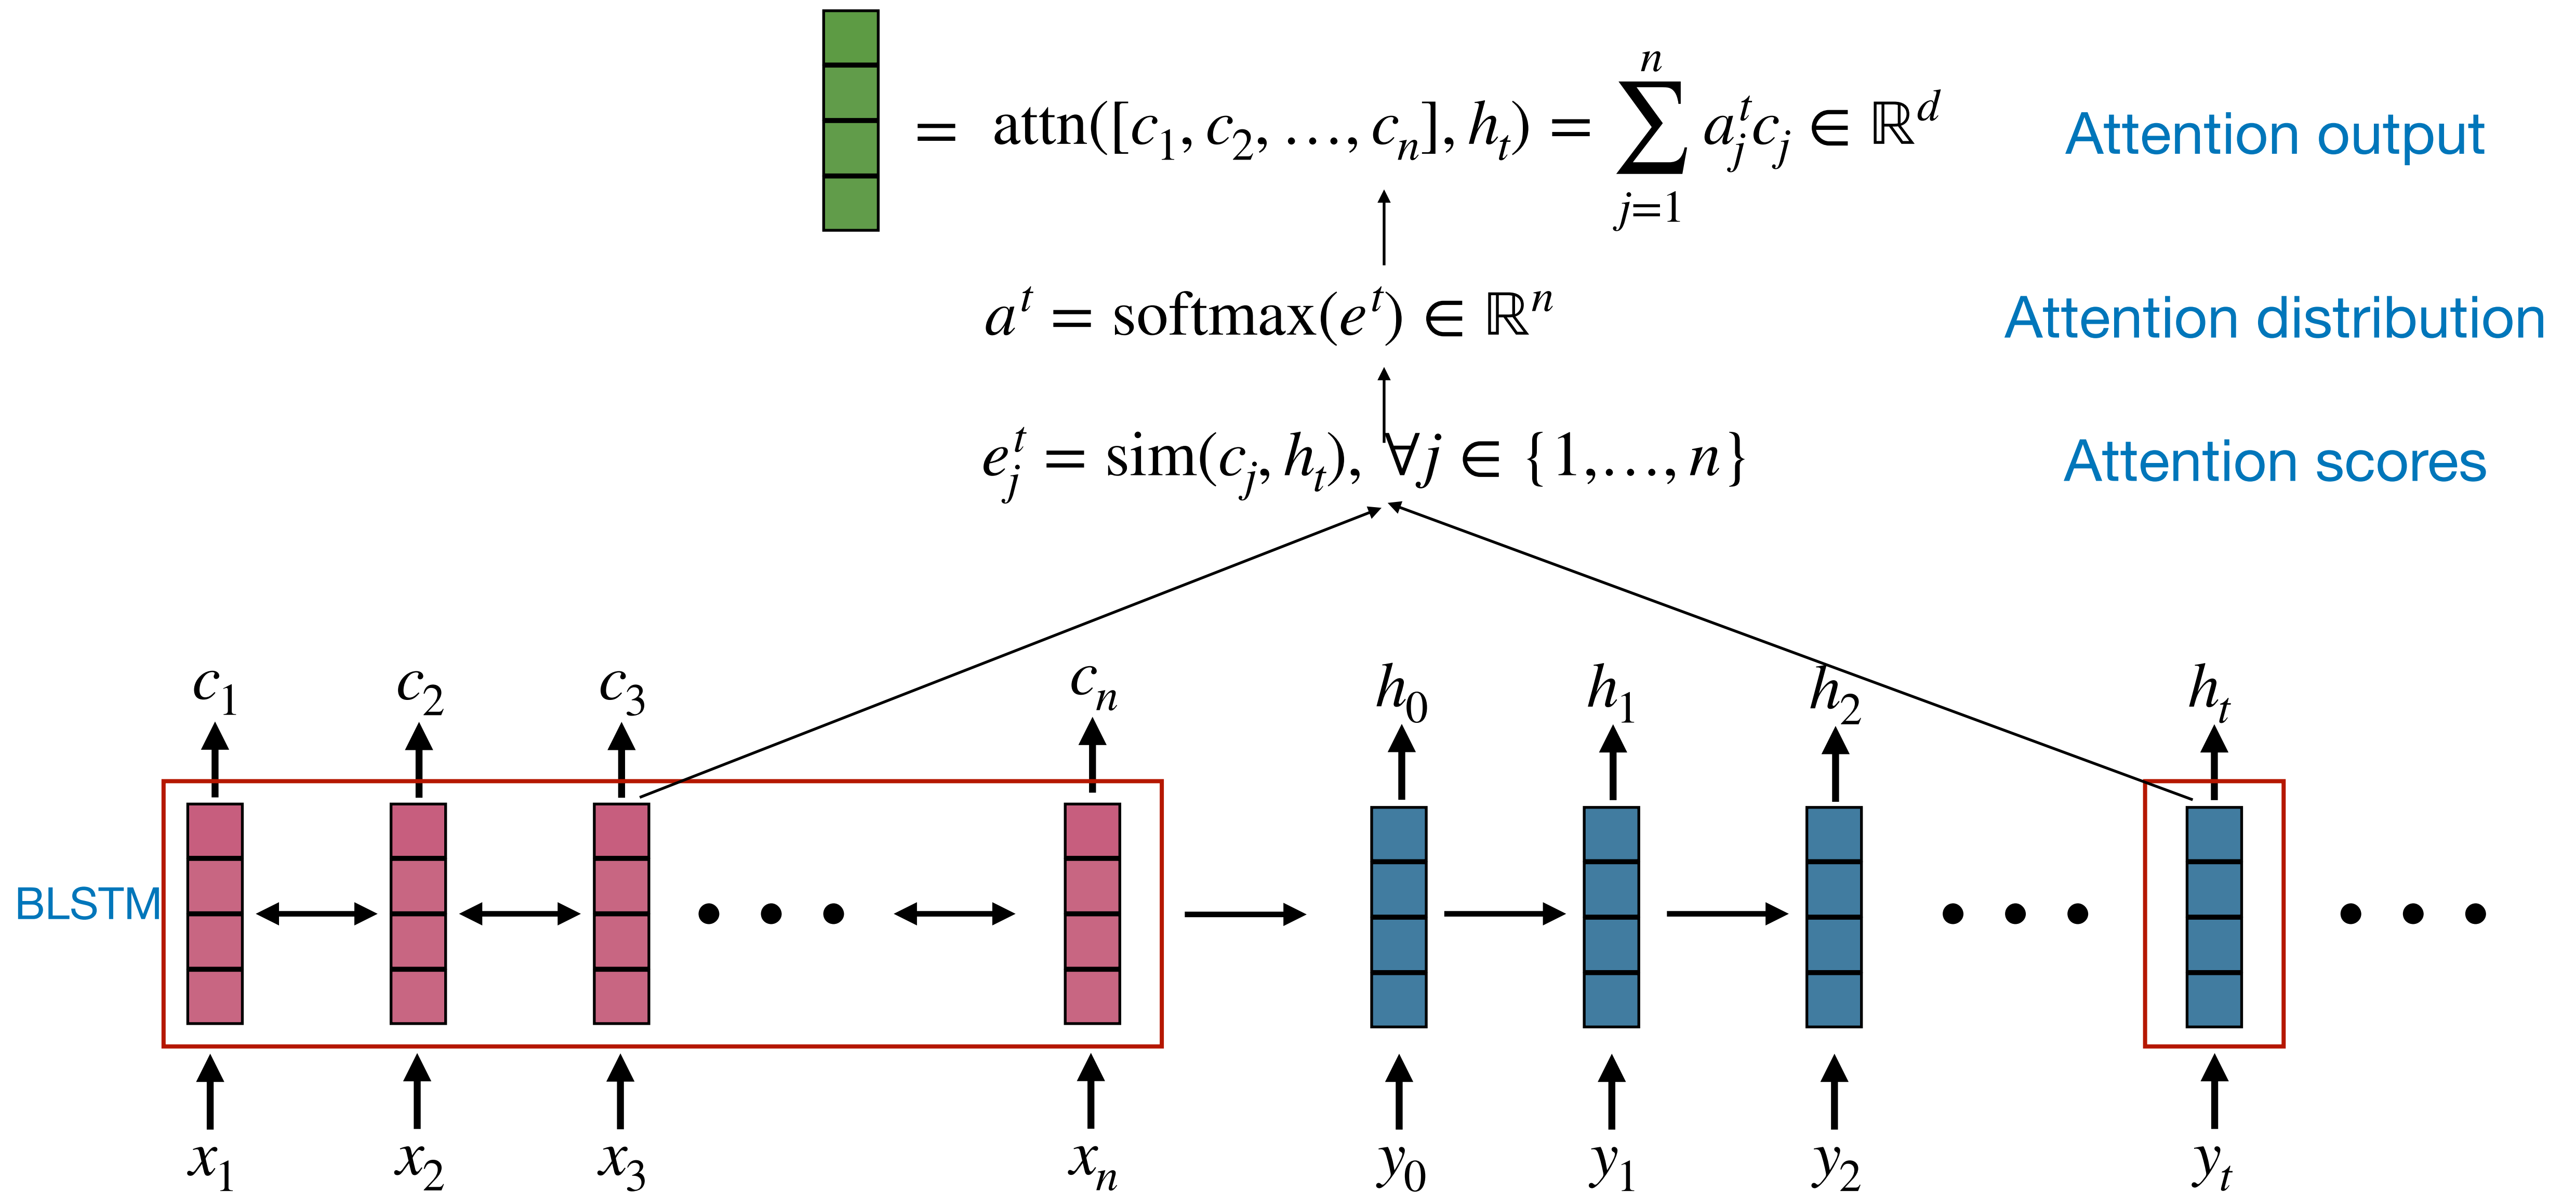


Revisit: Motivation of Attention Mechanism


- A single encoding vector needs to capture **all the information** about source sentence
- Longer sequences can lead to **vanishing gradients**




Recap: Attention Mechanism



Recap: MT Progress



MT Progress



Transformer

This Lecture

- Do we really need RNNs to model the arbitrary context?
- **Maybe attention is all you need!**

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin
illia.polosukhin@g


[CITATION] **Attention is all you need**

[A Vaswani](#) - Advances in Neural Information Processing Systems, 2017

☆ Save ↗ Cite Cited by 153082 Related articles ↗

Transformers

Transformer Encoder Block

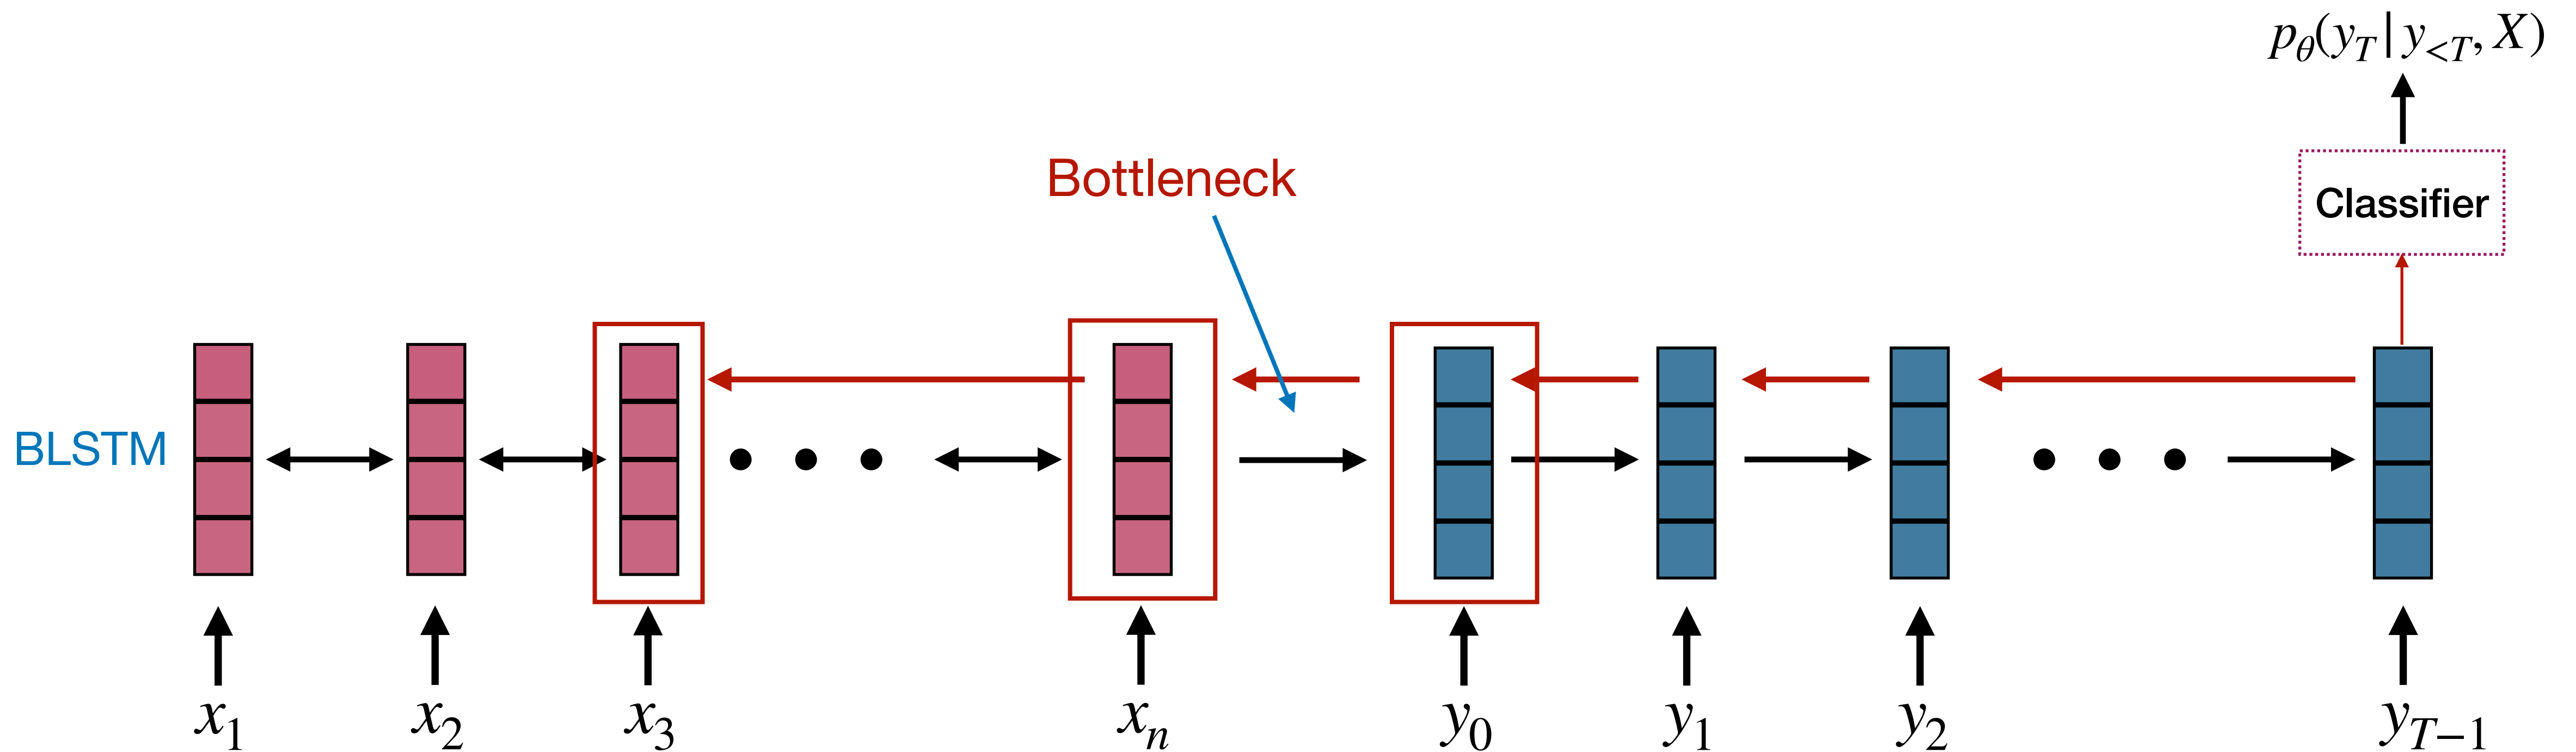


- **Three Key Components**
 - (Masked) Multi-head Self-Attention
 - Layer Normalization
 - Position-wise Feed-Forward Network

Self-Attention

Revisit: Motivation of Attention Mechanism

- A single encoding vector needs to capture **all the information** about source sentence
- Longer sequences can lead to **vanishing gradients**

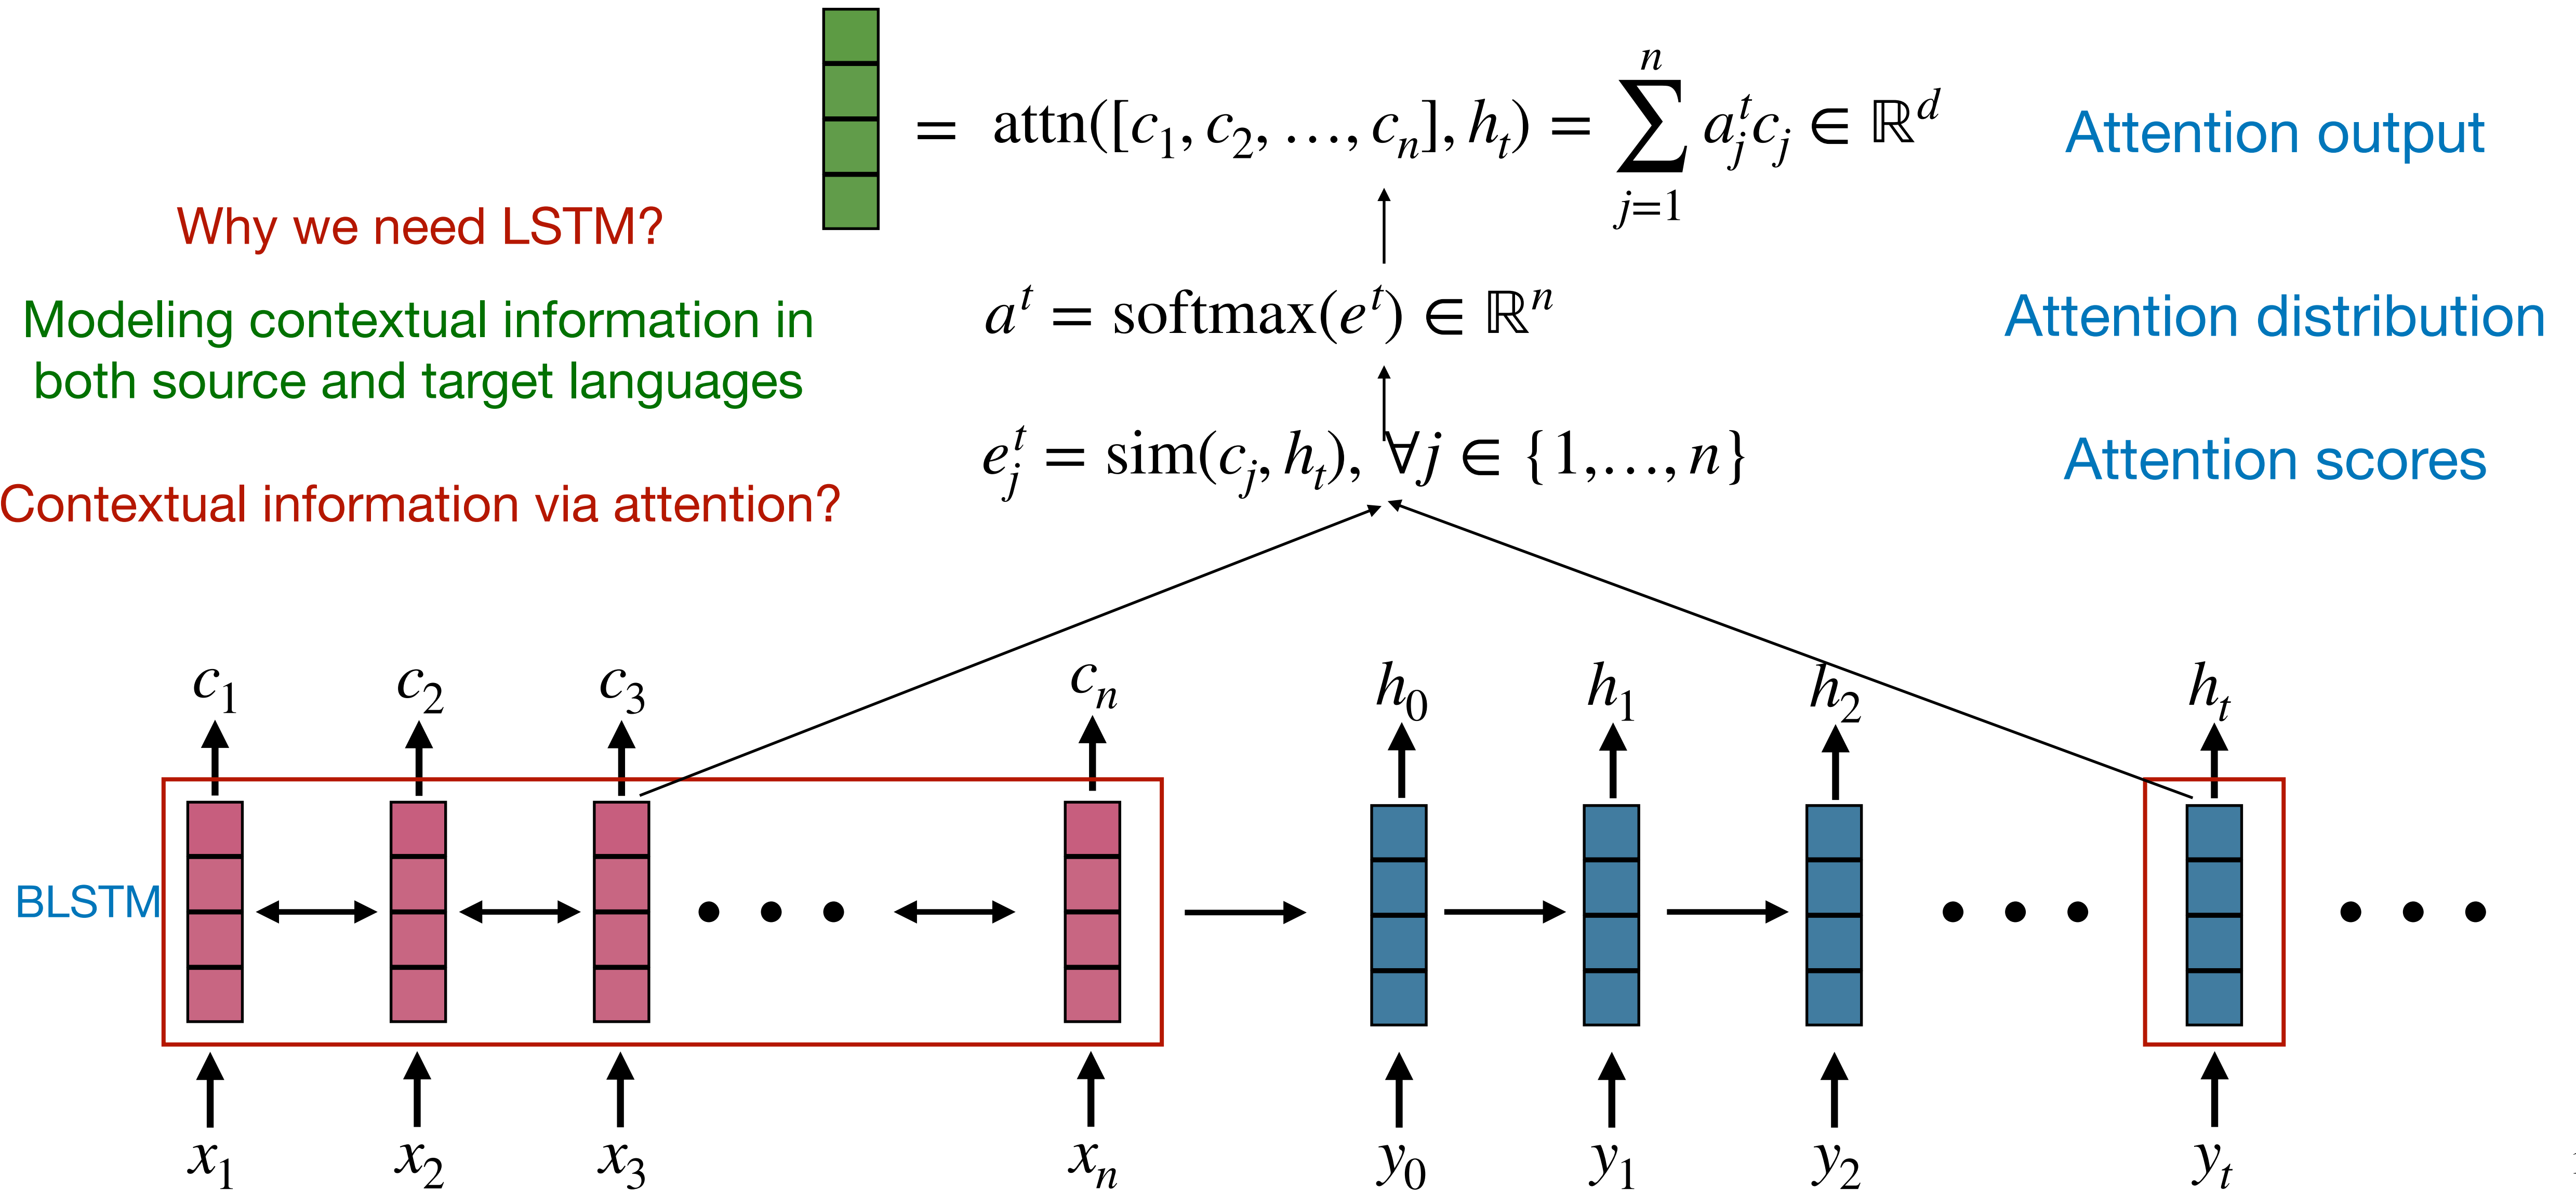


Recap: Attention Mechanism

Why we need LSTM?


Modeling contextual information in both source and target languages

Contextual information via attention?




Self-Attention

- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**




Self-Attention

- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**



Self-Attention

- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**





- Capturing long-distance dependencies
- No gradient vanishing
- Parallel computation!

Self-attention in equations

- A sequence of input vectors $x_1, \dots, x_n \in \mathbb{R}^d$
- First, construct a set of **queries**, **keys** and **values**:

$$q_i = W_Q x_i, \quad k_i = W_K x_i, \quad v_i = W_V x_i$$



Self-attention in equations

- A sequence of input vectors $x_1, \dots, x_n \in \mathbb{R}^d$
- First, construct a set of **queries**, **keys** and **values**:

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$

- Second, for each q_i , compute attention scores and attention distributions:

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right) \quad \text{aka. "scaled dot product"}$$

- Finally, compute the weighted sum:

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

Why *Scaled* Dot Product?

- **Softmax is sensitive to scale**

If $[x_1, x_2] = [0.1, 0.5]$, $\alpha = 10$

$$\text{softmax}([x_1, x_2]) = \left[\frac{e^{x_1}}{e^{x_1} + e^{x_2}}, \frac{e^{x_2}}{e^{x_1} + e^{x_2}} \right]$$

[0.4013, 0.5987]

$$\text{softmax}([\alpha x_1, \alpha x_2]) = \left[\frac{e^{\alpha x_1}}{e^{\alpha x_1} + e^{\alpha x_2}}, \frac{e^{\alpha x_2}}{e^{\alpha x_1} + e^{\alpha x_2}} \right]$$

[0.0180, 0.9820]

Self-attention in equations

- A sequence of input vectors $x_1, \dots, x_n \in \mathbb{R}^d$
- First, construct a set of **queries**, **keys** and **values**:

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$

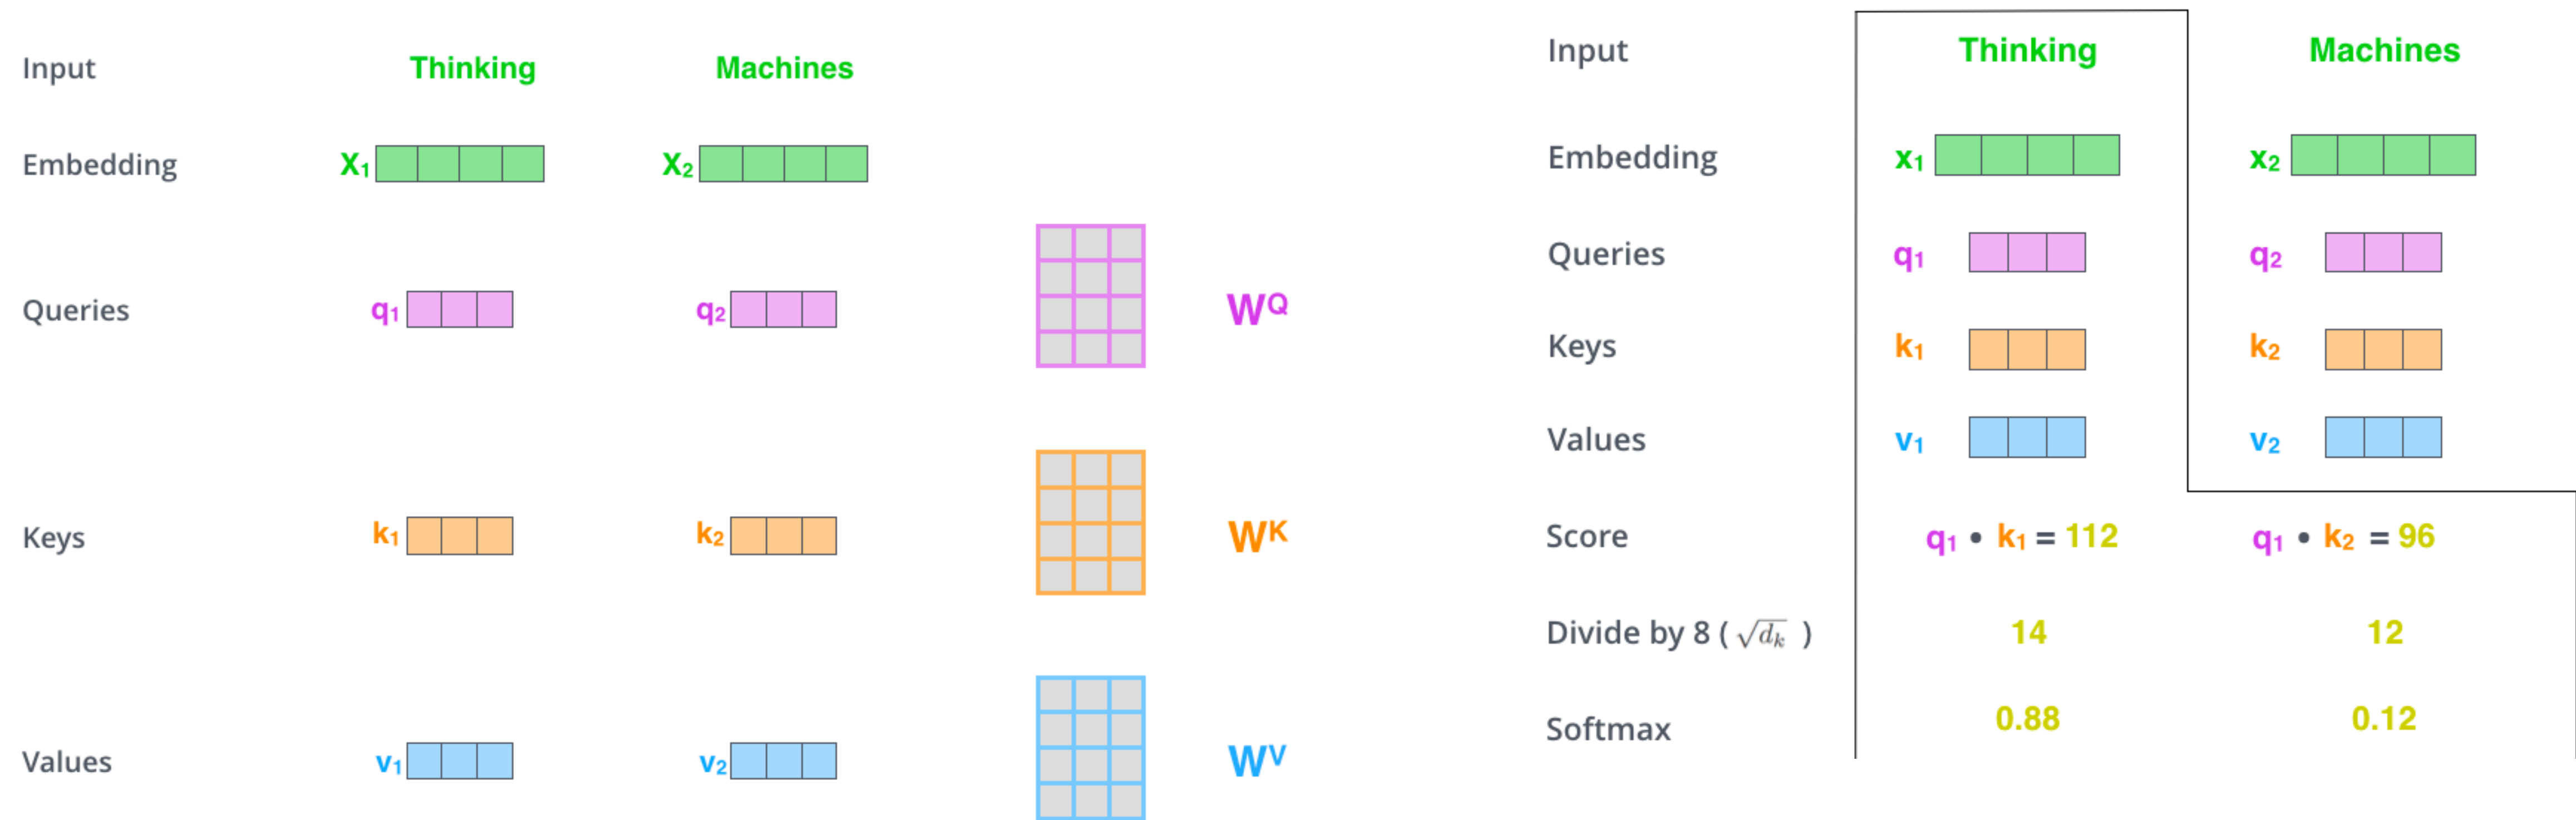
- Second, for each q_i , compute attention scores and attention distributions:

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right) \quad \text{aka. "scaled dot product"}$$

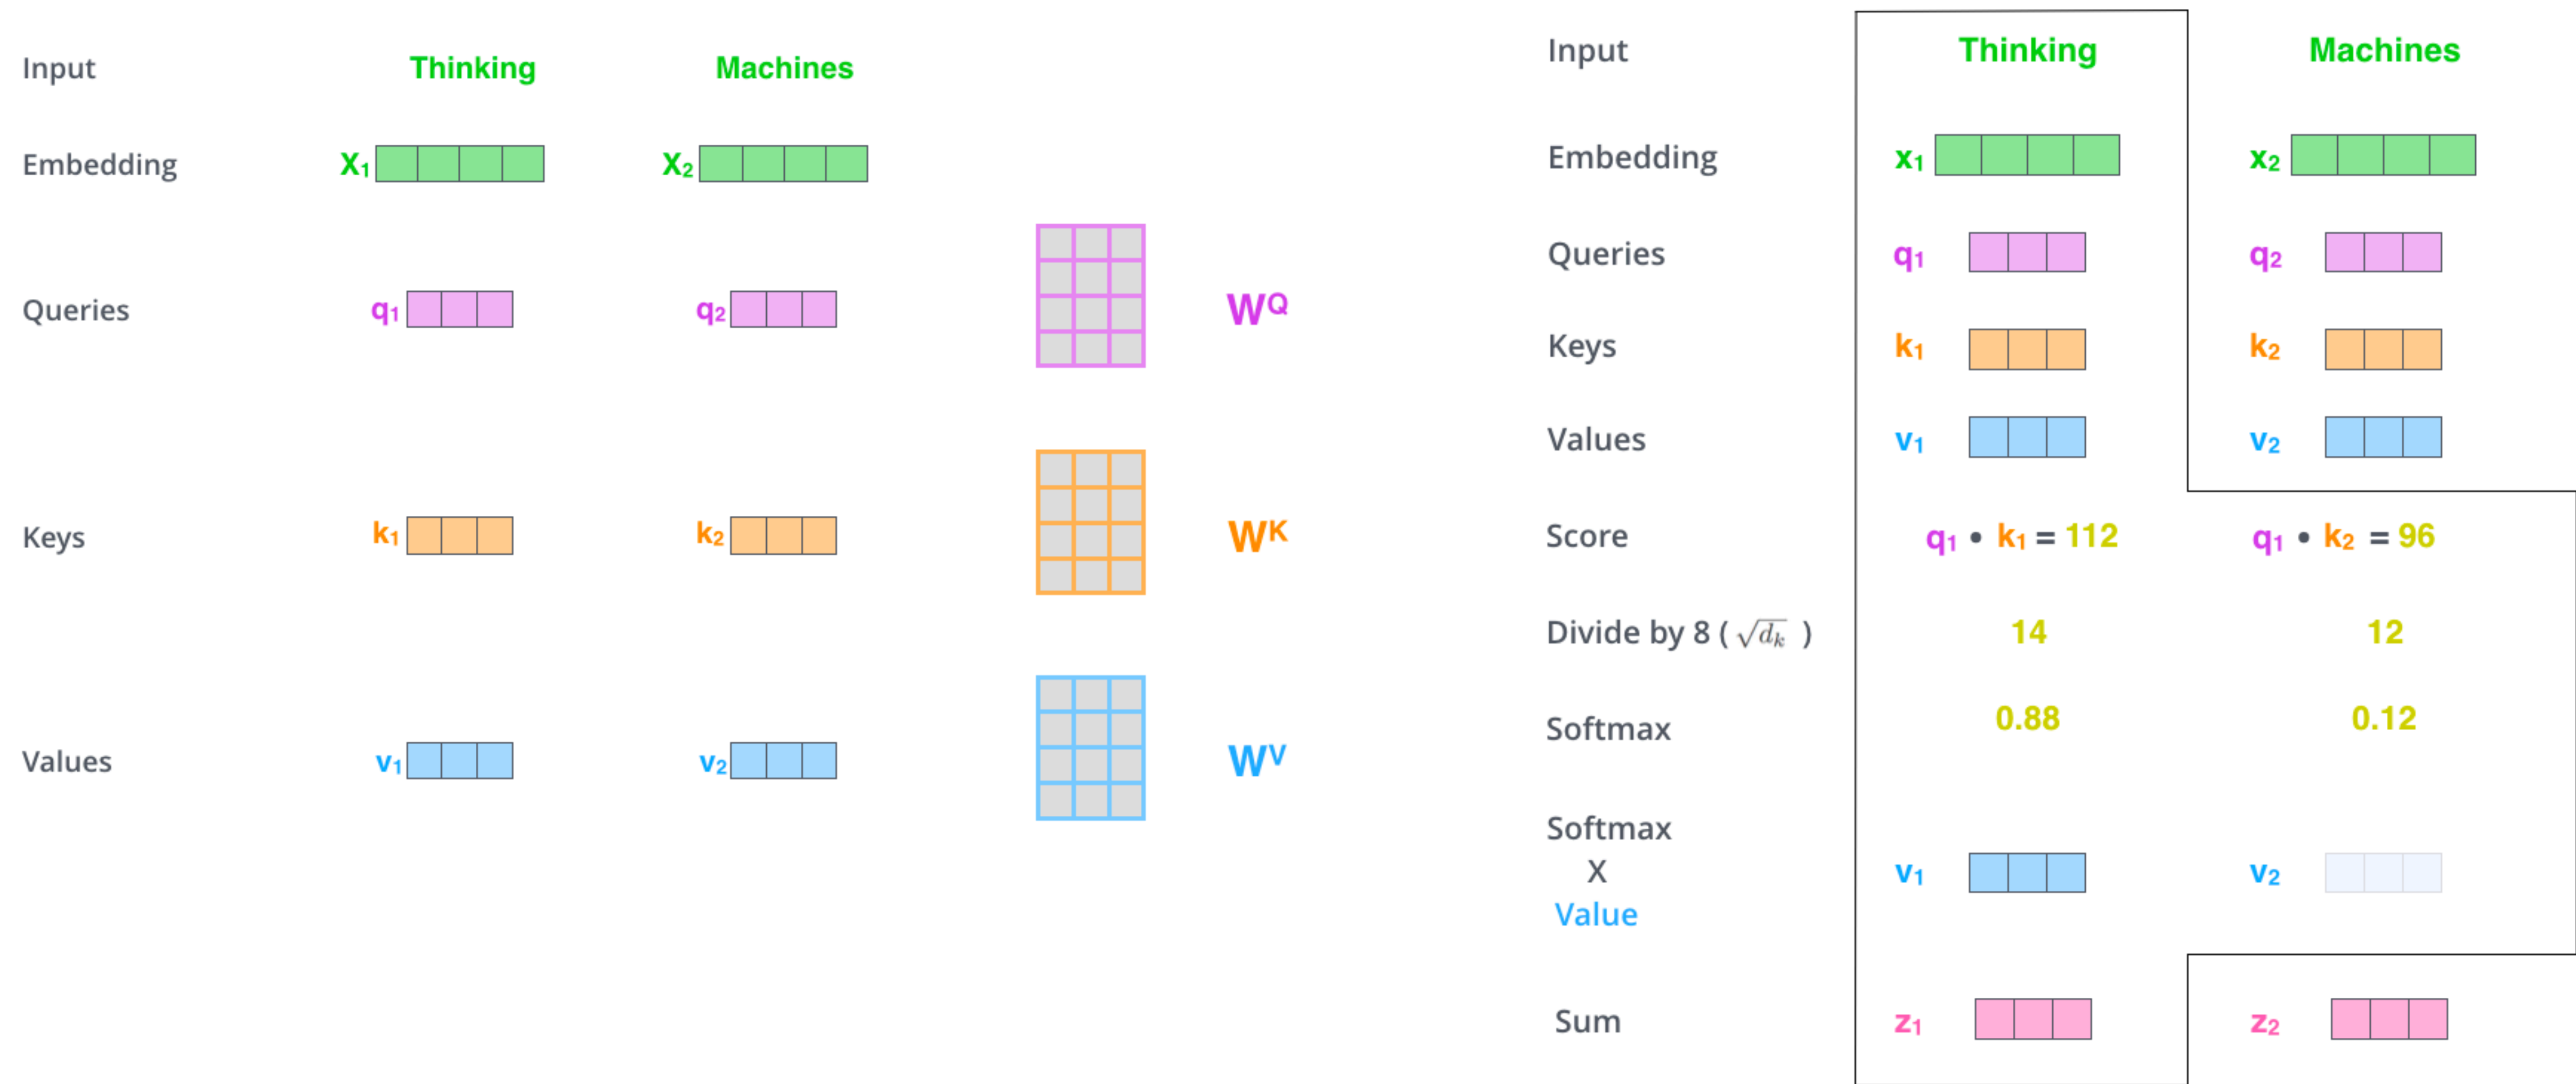
- Finally, compute the weighted sum:

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

Self-attention: Illustration



Self-attention: Illustration



<http://jalammar.github.io/illustrated-transformer/>

Self-attention: matrix notations

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

The diagram illustrates the self-attention mechanism using matrix notations and visual representations of matrices. It shows the calculation of the attention weights and the resulting output matrix.

The input matrices are:

- Q (Query matrix, represented by a 2x3 pink grid)
- K^T (Key matrix, represented by a 3x2 orange grid)
- V (Value matrix, represented by a 2x3 blue grid)

The attention mechanism is represented by the following expression:

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) V$$

The result of the softmax operation is a 2x3 pink grid, labeled Z .

The final output is the result of multiplying the attention weights (Z) by the value matrix (V), resulting in a 2x3 pink grid.

Self-attention: matrix notations



hardmaru
@hardmaru

The most important formula in deep learning after 2018

Self-Attention

What is self-attention? Self-attention calculates a weighted average of feature representations with the weight proportional to a similarity score between pairs of representations. Formally, an input sequence of n tokens of dimensions d , $X \in \mathbf{R}^{n \times d}$, is projected using three matrices $W_Q \in \mathbf{R}^{d \times d_q}$, $W_K \in \mathbf{R}^{d \times d_k}$, and $W_V \in \mathbf{R}^{d \times d_v}$ to extract feature representations Q , K , and V , referred to as query, key, and value respectively with $d_k = d_q$. The outputs Q , K , V are computed as

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V. \quad (1)$$

So, self-attention can be written as,

$$S = D(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_q}} \right) V, \quad (2)$$

where softmax denotes a *row-wise* softmax normalization function. Thus, each element in S depends on all other elements in the same row.

9:08 PM · Feb 9, 2021 · Twitter Web App

580 Retweets 38 Quote Tweets 3,407 Likes

Attention is *General*


- Given a set of **key** and **value** vectors, and a **query** vector, attention is a technique to compute a weighted sum of the **value** vectors, dependent on the **query and keys**
 - We sometimes say that the **query** attends to the **values** via **keys**
 - In the NMT case, each decoder hidden state (**query**) attends to all the encoder hidden states (**keys and values**)

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Attention is *General*

- Intuition

- The weighted sum is a *selective summary* of the information contained in the **values**, where the **query** and **keys** determines which **values** to focus on
- Attention is a way to obtain a *fixed-size representation* of an arbitrary set of representations (the **values**), dependent on some other representation (the **query**)



Multi-head Attention


- Problem with self-attention?

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

one set of attention weights a_i


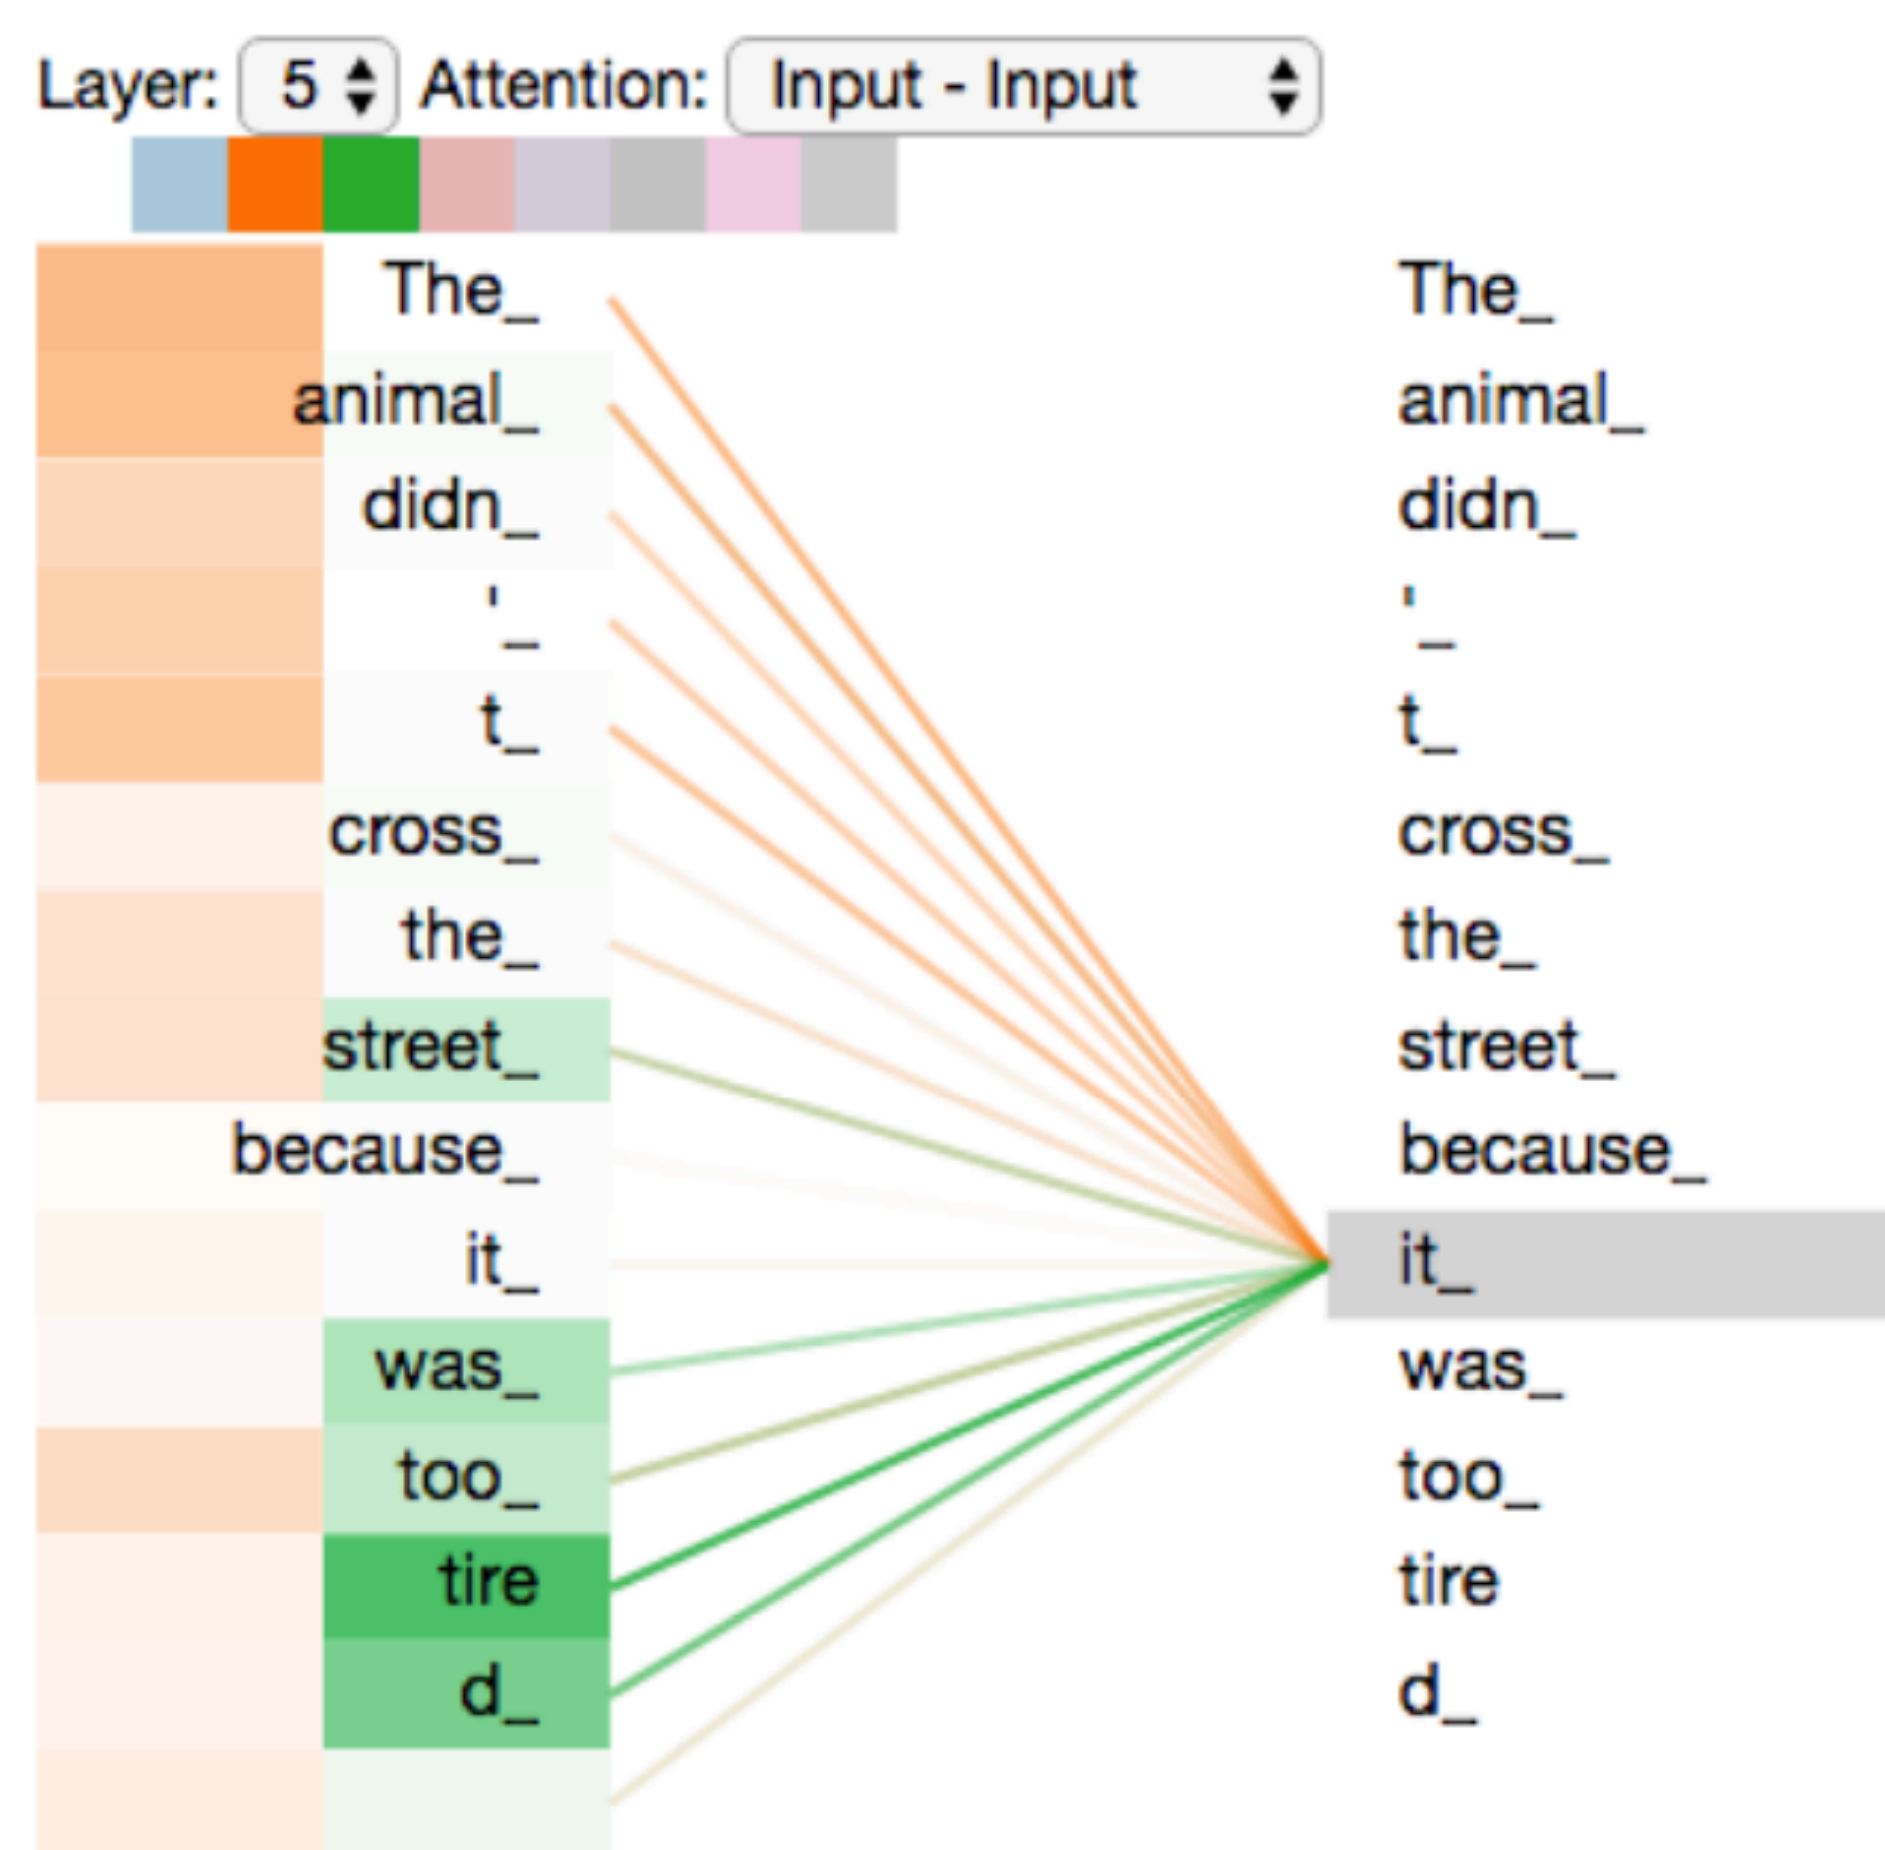
- It is better to use multiple attention weights instead of one!
 - Each attention can focus on different positions
- How to do this? Splits queries, keys, values to multiple heads!

Multi-head Attention: Head Split




$$h_1 = \text{attn}(Q_1, K_1, V_1) = \text{softmax}\left(\frac{Q_1 K_1^T}{\sqrt{d/2}}\right) V_1$$
$$h_2 = \text{attn}(Q_2, K_2, V_2) = \text{softmax}\left(\frac{Q_2 K_2^T}{\sqrt{d/2}}\right) V_2$$
$$Y = \text{concat}(h_1, h_2) W_O$$

What does multi-head attention learn?



Self-Attention

- **Self-attention:** attention within on single sequence
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**




- **How to apply to auto-regressive case?**

$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$

The diagram shows the equation with annotations. A red box labeled "Next Token" has a red arrow pointing to the y_t term in the denominator. A blue box labeled "history" has a blue arrow pointing to the $y_{<t}$ term in the denominator. The $y_{<t}$ term is also underlined in blue.


Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$




Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$



Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$




How to vectorize?

Masked Multi-Head Attention

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right)$$



Efficient implementation: compute attention as we normally do, mask out attention to future words by setting attention scores to $-\infty$


```
dot = torch.bmm(queries, keys.transpose(1, 2))

indices = torch.triu_indices(t, t, offset=1)
dot[:, indices[0], indices[1]] = float('-inf')


dot = F.softmax(dot, dim=2)
```

Missing Piece: Positional Information

RNN



Self-attention




Missing Piece: Positional Information

- Unlike RNNs, self-attention does **not** build in order information
 - Encode the order of the sentence into the input x_1, \dots, x_n
- Solution: add **positional encoding** to the input embeddings

$$x_i \leftarrow x_i + p_i$$

- Use sine and cosine functions of different frequencies (not learnable)

$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$



Transformer: Pros and Cons

- **Easier to capture dependencies:** we draw attention between every pair of words!
- **Easier to parallelize:**

$$\begin{aligned}\text{MultiHead}(X) &= \text{concat}(h_1, \dots, h_k)W_O \\ h_i &= \text{attn}(Q_i, K_i, V_i) \\ Q_i &= (XW_Q)^i, K_i = (XW_K)^i, V_i = (XW_V)^i\end{aligned}$$

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

- **Quadratic computation in self-attention:**
 - Can be very expensive when the sequence is very long: $O(hn^2 + nd)$
- **Harder to model positional information**

Transformer vs. RNN

RNN/LSTM

Transformer

- **Time Complexity**

$O(n)$

$O(n^2)$

- **Memory Complexity**

$O(n)$

$O(n^2)$

- **Training Speed**

Slow

Fast


- **Decoding Speed**

Fast

Slow

Transformers

Transformer Encoder Block



- **Three Key Components**

- (Masked) Multi-head Self-Attention
- Layer Normalization
- Position-wise Feed-Forward Network

Layer Normalization

Layer Normalization

- **Motivation:** normalize each vector individually to control vector scale


$$Y = \frac{X - E[X]}{\sqrt{\text{Var}[X] + \epsilon}} * \gamma + \beta$$

Initialization:

$$\gamma = 1, \beta = 0$$


What is the range of Y ?

$$(-\sqrt{d}, \sqrt{d})$$



Transformers

Transformer Encoder Block



- **Three Key Components**
 - (Masked) Multi-head Self-Attention
 - Layer Normalization
 - Position-wise Feed-Forward Network

Position-wise Feed Forward Network

Position-wise Feed Forward Network

- There is no elementwise nonlinearities in self-attention; stacking more self-attention layers just re-averages value vectors
- Simple fix: add a feed-forward network to post-process each output vector


$$\text{FFN}(\mathbf{x}_i) = W_2 \text{ReLU}(W_1 \mathbf{x}_i + \mathbf{b}_1) + \mathbf{b}_2$$

A large number
of parameters

$$W_1 \in \mathbb{R}^{d_{ff} \times d}, \mathbf{b}_1 \in \mathbb{R}^{d_{ff}}$$

$$W_2 \in \mathbb{R}^{d \times d_{ff}}, \mathbf{b}_2 \in \mathbb{R}^d$$

In practice, they use $d_{ff} = 4d$




Feed-Forward Layers

- Feed-forward layers constitute **two-thirds** of parameters
- Operates as memories of textual patterns (Gova et al., 2021)

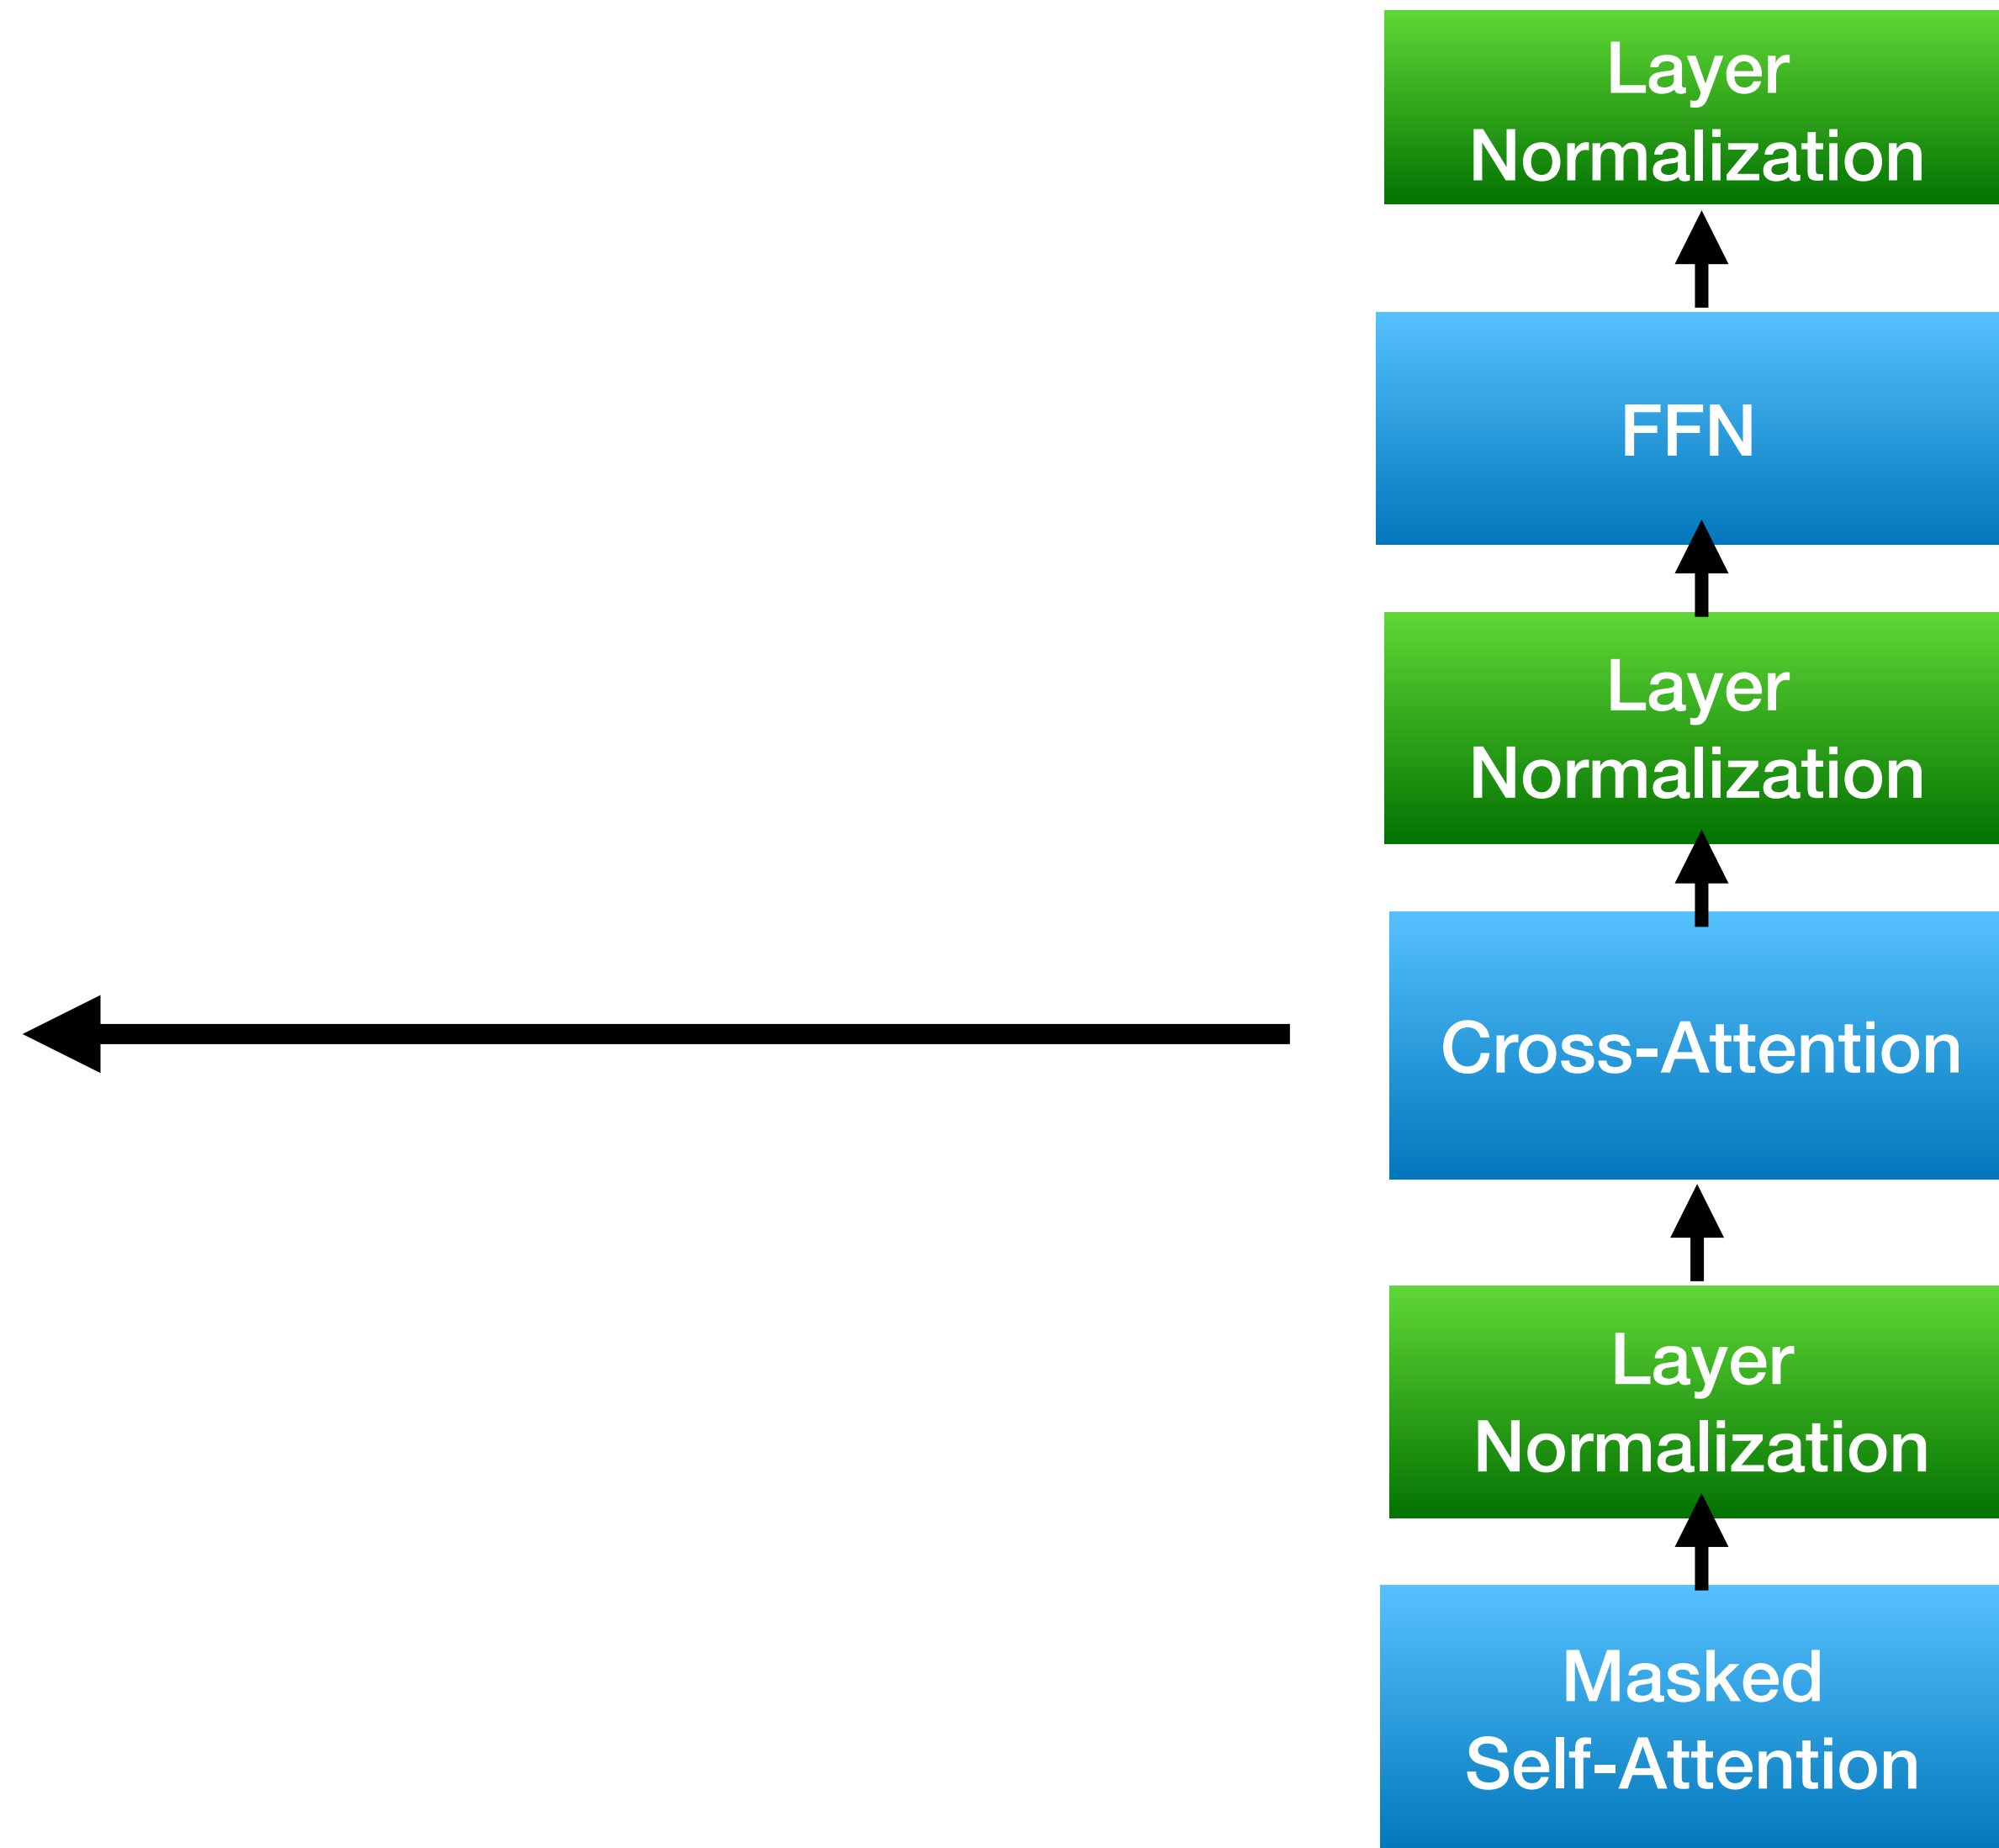
Key	Pattern	Example trigger prefixes
k_{449}^1	Ends with “ <i>substitutes</i> ” (shallow)	<i>At the meeting, Elton said that “for artistic reasons there could be no substitutes</i> <i>In German service, they were used as substitutes</i> <i>Two weeks later, he came off the substitutes</i>
k_{2546}^6	Military, ends with “ <i>base</i> ”/“ <i>bases</i> ” (shallow + semantic)	<i>On 1 April the SRSG authorised the SADF to leave their bases</i> <i>Aircraft from all four carriers attacked the Australian base</i> <i>Bombers flying missions to Rabaul and other Japanese bases</i>
k_{2997}^{10}	a “part of” relation (semantic)	<i>In June 2012 she was named as one of the team that competed</i> <i>He was also a part of the Indian delegation</i> <i>Toy Story is also among the top ten in the BFI list of the 50 films you should</i>
k_{2989}^{13}	Ends with a time range (semantic)	<i>Worldwide, most tornadoes occur in the late afternoon, between 3 pm and 7</i> <i>Weekend tolls are in effect from 7:00 pm Friday until</i> <i>The building is open to the public seven days a week, from 11:00 am to</i>
k_{1935}^{16}	TV shows (semantic)	<i>Time shifting viewing added 57 percent to the episode’s</i> <i>The first season set that the episode was included in was as part of the</i> <i>From the original NBC daytime version , archived</i>

Transformers

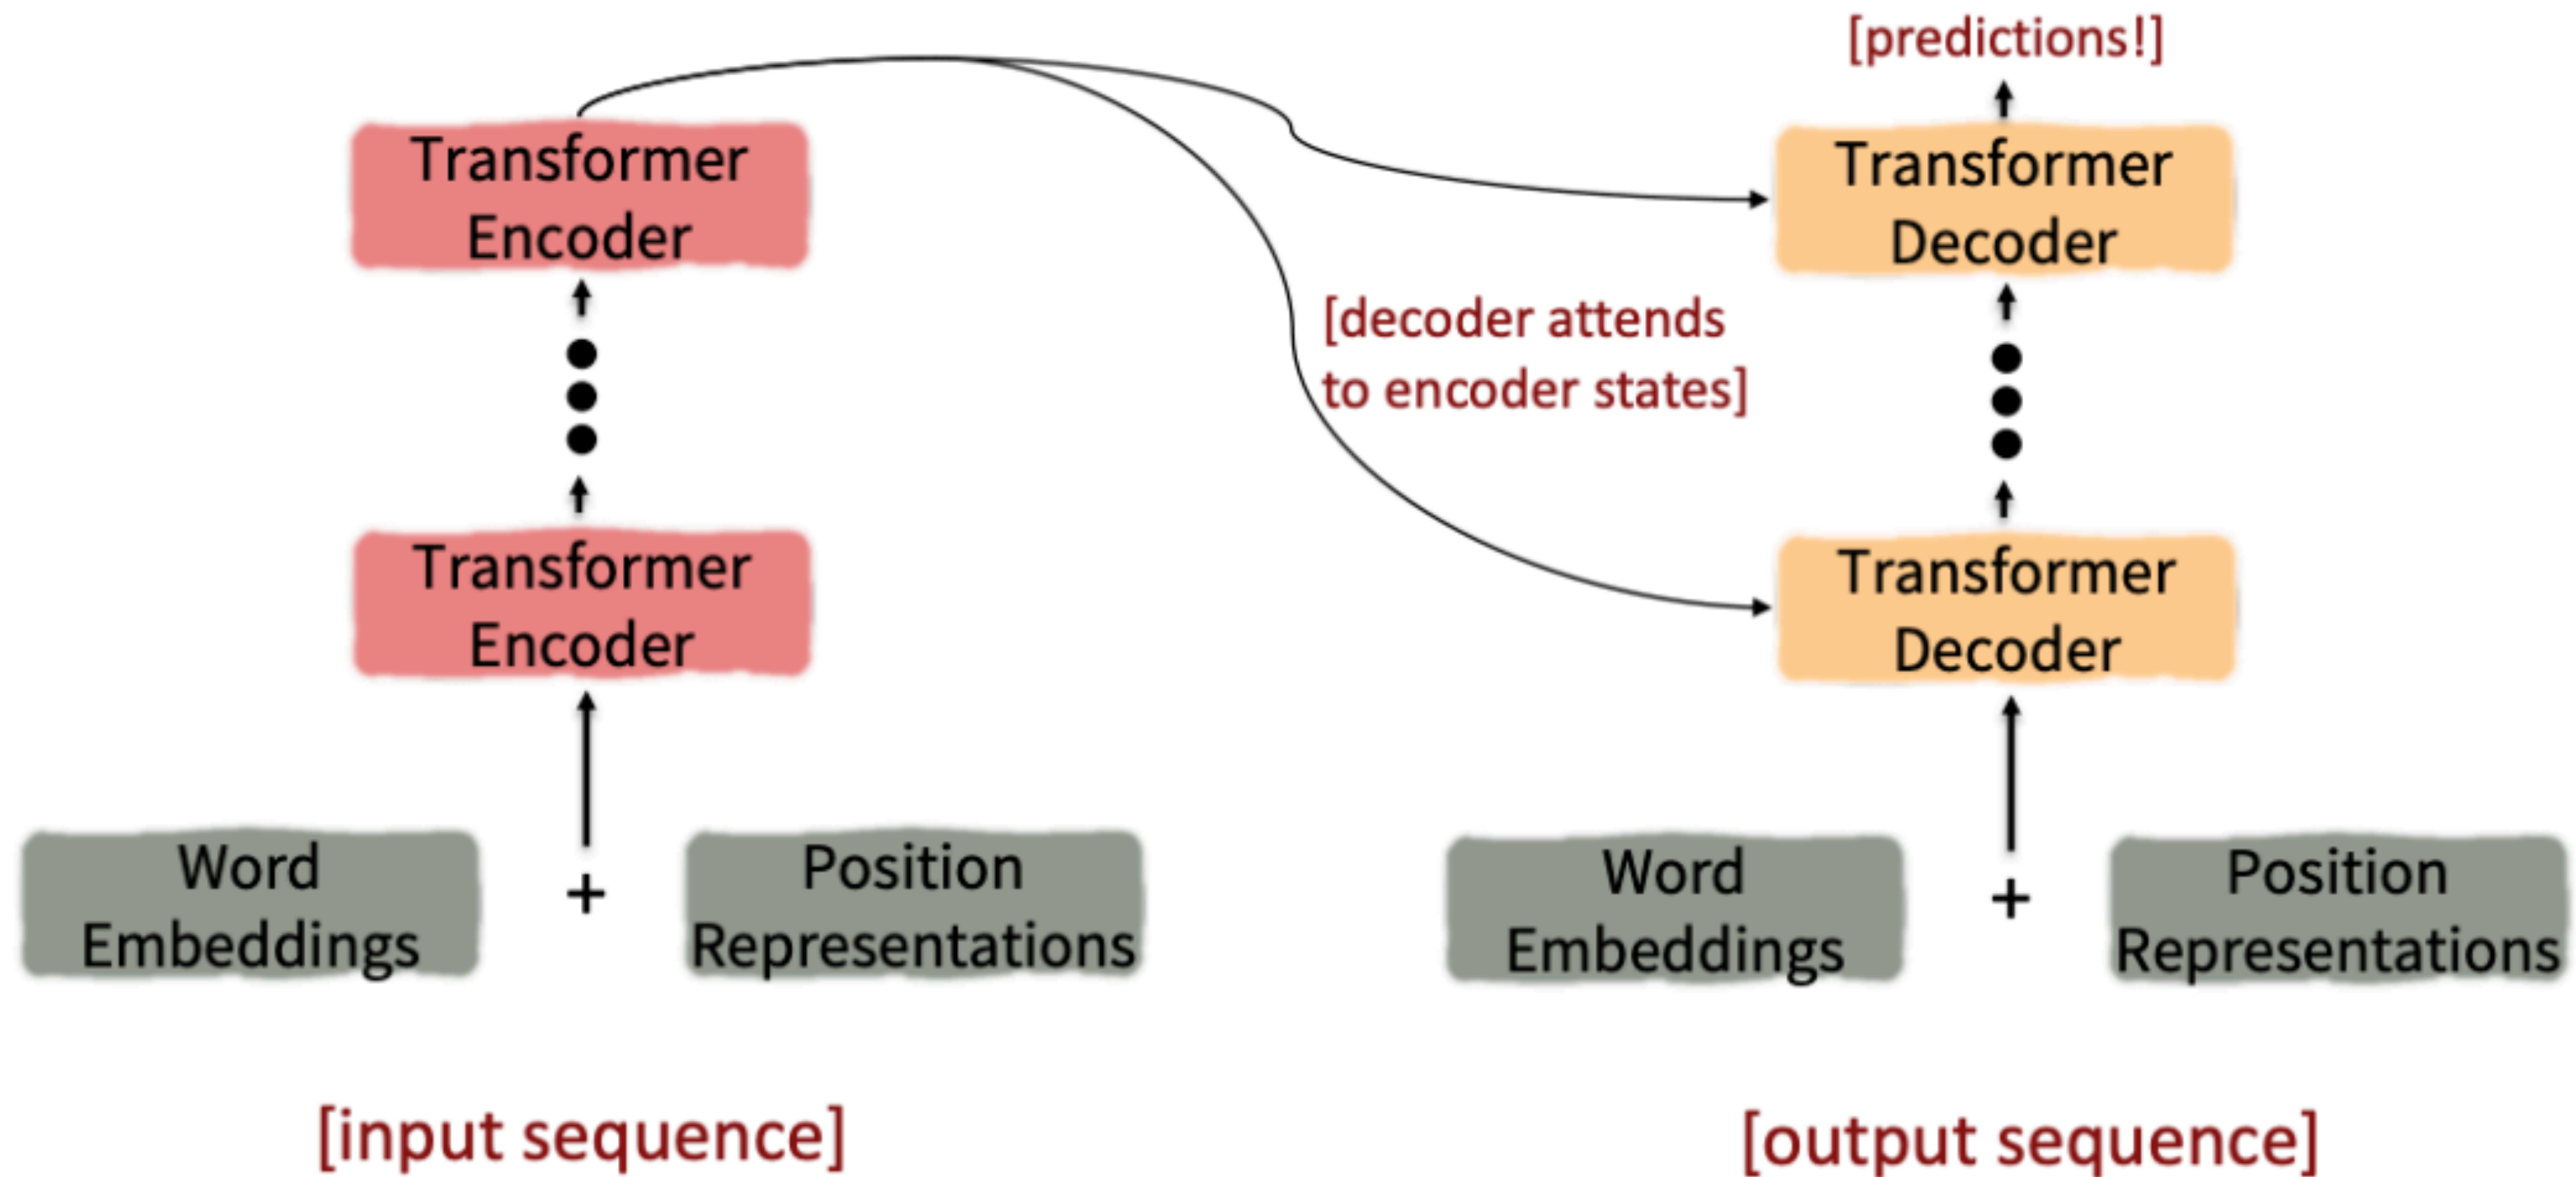
Transformer Encoder Block



Transformer Decoder Block



Putting the pieces together



Transformer: Machine Translation

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

Transformer: Document Generation

Model	Test perplexity	ROUGE-L
<i>seq2seq-attention, $L = 500$</i>	5.04952	12.7
<i>Transformer-ED, $L = 500$</i>	2.46645	34.2
<i>Transformer-D, $L = 4000$</i>	2.22216	33.6
<i>Transformer-DMCA, no MoE-layer, $L = 11000$</i>	2.05159	36.2
<i>Transformer-DMCA, MoE-128, $L = 11000$</i>	1.92871	37.9
<i>Transformer-DMCA, MoE-256, $L = 7500$</i>	1.90325	38.8

Significant gains compared to
seq2seq-attention with LSTMs